

# CE 440 Introduction to Operating System

## Lecture 9: Deadlock Fall 2025

Prof. Yigong Hu



Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

# Deadlock

## Synchronization is a live gun

- We can easily shoot ourselves in the foot
- Incorrect use of synchronization can block all processes
- You have likely been intuitively avoiding this situation already

# Example: Single-Lane Bridge Crossing



*CA 140 to Yosemite National Park*



# Deadlock

## Synchronization is a live gun

- We can easily shoot ourselves in the foot
- Incorrect use of synchronization can block all processes
- You have likely been intuitively avoiding this situation already

If one process tries to access a resource that a second process holds, and vice-versa, they can never make progress


We call this situation **deadlock**, and we'll look at:

- Definition and conditions necessary for deadlock
- Representation of deadlock conditions
- Approaches to dealing with deadlock

# Bridge Crossing Example

**Each segment of road can be viewed as a resource**

- Car must own the segment under them
- Must acquire segment that they are moving into



**Deadlock resolved if one car backs up (preempt resources and rollback)**

**Starvation: East-going traffic really fast → no one gets to go west**

# Deadlock Definition

**Deadlock is a problem that can arise:**

- When processes compete for access to limited resources
- When processes are incorrectly synchronized

**Definition:**

- Deadlock exists among a set of processes if every process is waiting for an event that can be caused only by another process in the set.

# Deadlock Example


```
mutex_t x, y;
```

Thread A:

```
void p1(void *ignored) {  
    lock(x);  
    lock(y);  
    /* critical section */  
    unlock(y);  
    unlock(x);  
}
```

Thread B:

```
void p2(void *ignored) {  
    lock(y);  
    lock(x);  
    /* critical section */  
    unlock(x);  
    unlock(y);  
}
```



# Deadlock Example: “Unlucky” Case


```
mutex_t x, y;
```

Thread A:

```
void p1(void *ignored) {  
    lock(x);  
  
    lock(y); stalled  
    <unreachable>  
    /* critical section */  
    unlock(y);  
    unlock(x);  
}
```

Thread B:

```
void p2(void *ignored) {  
  
    lock(y);  
  
    lock(x); stalled  
    <unreachable>  
    /* critical section */  
    unlock(x);  
    unlock(y);  
}
```



Neither thread will get to run → Deadlock

# Deadlock Example: “Lucky” Case


```
mutex_t m1, m2;
```

Thread A:

```
void p1(void *ignored) {  
    lock(x);  
    lock(y);  
  
    /* critical section */  
    unlock(y);  
    unlock(x);  
}
```

Thread B:

```
void p2(void *ignored) {  
  
    lock(y);  
    lock(x);  
  
    /* critical section */  
    unlock(x);  
    unlock(y);  
}
```



Sometimes, schedule won't trigger deadlock!

# Deadlock Example


```
mutex_t m1, m2;
```

Thread A:

```
void p1(void *ignored) {  
    lock(x);  
    lock(y);  
    /* critical section */  
    unlock(y);  
    unlock(x);  
}
```

Thread B:

```
void p2(void *ignored) {  
    lock(y);  
    lock(x);  
    /* critical section */  
    unlock(x);  
    unlock(y);  
}
```



**This lock pattern exhibits non-deterministic deadlock**

- Sometimes it happens, sometimes it doesn't!

**This is really hard to debug!**

# **Deadlock Questions**

**Can you have deadlock w/o mutexes?**

# Deadlock Example: Memory Contention

```
mutex_t m1, m2;
```

Thread A:

```
void p1(void *ignored) {  
    AllocateOrWait(1 MB)  
    AllocateOrWait(1 MB)  
    /* do something */  
    Free(m2);  
    unlock(m1);  
}
```

Thread B:

```
void p2(void *ignored) {  
    AllocateOrWait(1 MB)  
    AllocateOrWait(1 MB)  
    /* critical section */  
    unlock(m1);  
    unlock(m2);  
}
```

If only 2 MB of space, we get same deadlock situation

# Deadlock Example: Memory Contention

```
mutex_t m1, m2;
```

Thread A:

```
void p1(void *ignored) {  
    AllocateOrWait(1 MB)  
  
    AllocateOrWait(1 MB)  
    /* do something */  
    Free(m2);  
    unlock(m1);  
}
```

Thread B:

```
void p2(void *ignored) {  
  
    AllocateOrWait(1 MB)  
  
    AllocateOrWait(1 MB)  
    /* critical section */  
    unlock(m1);  
    unlock(m2);  
}
```


If only 2 MB of space, we get same deadlock situation

# Deadlock Questions

## Can you have deadlock w/o mutexes?

- Threads often block waiting for resources
  - Locks
  - Terminals
  - Printers
  - CD drives
  - Memory
- Same problem with condition variables
  - Suppose resource 1 managed by c1, resource 2 by c2
  - A has 1, waits on c2, B has 2, waits on c1
- Threads often block waiting for other threads
  - Pipes
  - Sockets
- You can deadlock on any of these!

# Dining Philosophers Problem




**Philosophers spend their lives  
alternating thinking and eating**

**Don't interact with neighbors,  
occasionally eat**

- Need 2 chopsticks to eat
- Release both when done

**Can only pick up 1 fork at a time**



# Dining Philosophers in Code

```
#define N 5 /* number of philosophers */
semaphore forks[N]; /* semaphores for each fork, each initialized to 1 (omitted) */
```

```
void philosopher(int i) /* i: philosopher id, 0 to 4 */
{
    while (true) {
        think(); /* philosopher is thinking */
        take_fork(i); /* take left fork */
        take_fork((i + 1) % N); /* take right fork */
        eat(); /* yum-yum, spaghetti */
        put_fork(i); /*put left fork back on the table*/
        put_fork((i + 1) % N); /* put right fork back on
the table */
    }
}
```

```
void take_fork(int i) {
    forks[i].P();
    /*wait for ith fork's semaphore*/
}

void put_fork(int i) {
    forks[i].V();
    /*signal ith fork's semaphore*/
}
```

What problem with this code?

# How to Avoid Deadlock Here?

**Multiple solutions exist:**

# How to Avoid Deadlock Here?

**Multiple solutions exist:**

**Simple one: allow at most 4 philosophers to sit simultaneously at the table**

**Another solution: define a partial order for resources (forks)**

- Number the forks
- Philosopher must always pick up lower-numbered fork first and then higher-numbered fork
- What happens if four philosophers all pick up their lower-numbered fork?
- Disadvantage
  - Not always practical, when the complete list of all resources is not known in advance

**Third solution: all or none each time**

# 2<sup>nd</sup> Attempt to Dining Philosopher Problem

Fix the previous code

```
#define N 5 /* number of philosophers */
semaphore forks[N]; /* semaphores for each fork, each initialized to 1 (omitted) */
```

```
void philosopher(int i) /* i: philosopher id, 0 to 4
*/
{
    while (true) {
        think(); /* philosopher is thinking */
        take_fork(i); /* take left fork */
        take_fork((i + 1) % N); /* take right fork */
        eat(); /* yum-yum, spaghetti */
        put_fork(i); /*put left fork back on the table*/
        put_fork((i + 1) % N); /* put right fork back on
the table */
    }
}
```

```
void take_fork(int i) {
    forks[i].P();
    /*wait for ith fork's semaphore*/
}

void put_fork(int i) {
    forks[i].V();
    /*signal ith fork's semaphore*/
}
```

# 2<sup>nd</sup> Attempt to Dining Philosopher Problem

```
#define N 5 /* number of philosophers */
#define LEFT (i+N-1) % N /* i's left neighbor */
#define RIGHT (i+1) % N /* i's right neighbor */
enum State {THINKING, HUNGRY, EATING}; /* a philosopher's status */
enum State states[N]; /* keep track of each philosopher's status */
semaphore mutex = 1; /* mutual exclusion for critical section */
semaphore phis[N]; /* semaphore for each philosopher, init to 0 */

void philosopher(int i) /* i: philosopher id, 0 to N-1 */
{
    while (true) {
        think(); /* philosopher is thinking */
        take_forks(i); /* take both forks */
        eat(); /* yum-yum, spaghetti */
        put_forks(i); /* put both forks */
    }
}
```

# 2<sup>nd</sup> Attempt to Dining Philosopher Problem

```
void take_forks(int i) /* i: philosopher id, 0 to N-1 */  
{  
    mutex.P(); /* enter critical section */  
    states[i] = HUNGRY; /* indicate philosopher is hungry */  
    test(i); /* try to acquire two forks */  
    mutex.V(); /* exit critical section */  
    phis[i].P(); /* block if forks not acquired */  
}  
  
void put_forks(int i) { /* i: philosopher id, 0 to N-1 */  
    mutex.P(); /* enter critical section */  
    states[i] = THINKING; /* indicate i finished eating */  
    test(LEFT); /* see if left neighbor can eat now */  
    test(RIGHT); /* see if right neighbor can eat now */  
    mutex.V(); /* exit critical section */  
}
```

# 2<sup>nd</sup> Attempt to Dining Philosopher Problem

```
void test(int i) /* i: philosopher id, 0 to N-1 */  
{  
    if (states[i] == HUNGRY &&  
        states[LEFT] != EATING &&  
        states[RIGHT] != EATING) {  
        states[i] = EATING; /* philosopher I can eat now */  
        phis[i].V(); /* signal i to proceed */  
    }  
}
```

# Notes for the Solution

**What is the purpose of states array?**

- given that already have the semaphore array?
- A semaphore doesn't have operations for checking its value!

**What if we don't use the mutex semaphore?**

**Why the semaphore array is for each philosopher?**

- Our first attempt uses semaphore array for each fork

**What if we put `phis[i].P();` inside the critical section?**

**What if we don't call the two test in `put_forks`?**

# Conditions for Deadlock

- Mutual exclusion – At least one resource must be held in a non sharable mode
- Hold and wait – There must be one process holding one resource and waiting for another resource
- No preemption – Resources cannot be preempted (critical sections cannot be aborted externally)
- Circular wait – There must exist a set of processes [P1, P2, P3,...,Pn] such that P1 is waiting for P2, P2 for P3, etc.

# Questions

**How to detect deadlocks?**


# Conditions for Deadlock

## View system as graph


- Processes and Resources are nodes
- Resource Requests and Assignments are edges

## Resource-Allocation Graph:


- A set of Threads  $T_1, T_2, \dots, T_n$
- Resource types  $R_1, R_2, \dots, R_n$ 
  - CPU cycles, memory space, I/O devices
- Each resource type  $R_i$  has  $W_i$  instances
- Each thread utilizes a resource as follows:
  - Request() / Use() / Release()
- Thread  $T_i$  requesting resource  $R_j$ :
- Thread  $T_i$  holding instance of  $R_j$ :




# Resource-Allocation Graph Example



# Resource-Allocation Graph Example



# Is This Deadlock?



# Deadlock Detection

**If graph has no cycles  $\Rightarrow$  no deadlock**

**If graph contains a cycle**

- Definitely deadlock if only one instance per resource (waits-for graph (WFG))
- Otherwise, maybe deadlock, maybe not

**Traverse the resource graph is expensive**

- Many processes and resources to traverse

**Only invoke detection algorithm periodically**

# Deal with Deadlock

**There are four approaches for dealing with deadlock:**

- Ignore it
- Prevention: write your code to make it impossible for deadlock to happen
- Avoidance – control allocation of resources
- Recovery – look for a cycle in dependencies

# Prevent by Eliminating One Condition

## 1. Mutual exclusion

- Buy more resources, split into pieces, or virtualize to make "infinite" copies
- Give illusion of infinite resources (e.g. virtual memory)

# Virtually Infinite Resources

Thread A:

```
void p1(void *ignored) {  
    AllocateOrWait(1 MB)  
    AllocateOrWait(1 MB)  
    /* do something */  
    Free(m2);  
    unlock(m1);  
}
```

Thread B:

```
void p2(void *ignored) {  
    AllocateOrWait(1 MB)  
    AllocateOrWait(1 MB)  
    /* critical section */  
    unlock(m1);  
    unlock(m2);  
}
```

With virtual memory we have “infinite” space so everything will just succeed, thus above example won’t deadlock

# Prevent by Eliminating One Condition

## 1. Mutual exclusion

- Buy more resources, split into pieces, or virtualize to make "infinite" copies
- Give illusion of infinite resources (e.g. virtual memory)

## 2. Hold and wait

- Wait on all resources at once (must know in advance)


## 3. No preemption

- Physical memory: virtualized with VM, can take physical page away and give to any process!

## 4. Circular wait

- Partial ordering of resources
  - e.g., always acquire mutex  $m_1$  before  $m_2$
  - Usually design locking discipline for application this way

# Request Resource in Partial Order



Thread A:

```
void p1(void *ignored) {  
    lock(x);  
    lock(y);  
    /* critical section */  
    unlock(y);  
    unlock(x);  
}
```

Thread B:

```
void p2(void *ignored) {  
    lock(y);  
    lock(x);  
    /* critical section */  
    unlock(x);  
    unlock(y);  
}
```

```
mutex_t x, y;
```



Thread A:

```
void p1(void *ignored) {  
    lock(x);  
    lock(y);  
    /* critical section */  
    unlock(y);  
    unlock(x);  
}
```

Thread B:

```
void p2(void *ignored) {  
    lock(x);  
    lock(y);  
    /* critical section */  
    unlock(x);  
    unlock(y);  
}
```

# Prevent by Eliminating One Condition

## 4. Circular wait

- Partial ordering of resources
  - e.g., always acquire mutex  $m_1$  before  $m_2$
  - Usually design locking discipline for application this way
- Make all threads request everything they'll need at the beginning.
  - Problem: Predicting future is hard, tend to over-estimate resources

# Recovering from Deadlock

## Terminate processes

- Abort all deadlocked processes
  - Processes need to start over again
- Abort one process at a time until cycle is eliminated
  - System needs to rerun detection after each abort

## Preempt resources (force their release)

- Need to select process and resource to preempt
- Need to rollback process to previous state
- Need to prevent starvation

## Roll back actions of deadlocked threads

- Common technique in databases (transactions)

# Avoid Deadlock

**Idea solution: When a process requests a resource, OS only grant it when:**

- The process can obtain all resources it needs in future requests
- Information in advance about what resources will be needed by processes to guarantee that deadlock will not happen

**Tough**

- Hard to determine all resources needed in advance
- Good theoretical problem, not as practical to use

# Three States

## Safe state

- System can delay resource acquisition to prevent deadlock

## Unsafe state

- No deadlock yet...
  - But threads can request resources in a pattern that unavoidably leads to deadlock
- Deadlock avoidance: prevent system from reaching an *unsafe* state

## Deadlocked state

- There exists a deadlock in the system
- Also considered “unsafe”

# Banker's Algorithm

- 1. Each process must state its maximum resource demand**
  - OS tracks available resource, maximum demand of each process
  
- 2. When a process requests resources:**
  - OS check whether the request would lead to an unsafe state

10 units of resource A

P1: Max = 7, Allocated = 3

P2: Max = 5, Allocated = 2

P3: Max = 3, Allocated = 2

# Deadlock Summary

**Deadlock occurs when processes are waiting on each other and cannot make progress**

- Cycles in Resource Allocation Graph (RAG)

**Deadlock requires four conditions**

- Mutual exclusion, hold and wait, no resource preemption, circular wait

**Four approaches to dealing with deadlock:**

- Ignore it – Living life on the edge
- Prevention – Make one of the four conditions impossible
- Avoidance – Banker's Algorithm (control allocation)
- Detection and Recovery – Look for a cycle, preempt or abort

# Condition Vars & Locks

C/Vs are also used without monitors in conjunction with locks

- void cond\_init (cond\_t \*, ...);
- void cond\_wait (cond\_t \*c, mutex\_t \*m);
  - Atomically unlock mand sleep until csignaled
  - Then re-acquire m and resume executing
- void cond\_signal (cond\_t \*c);
- void cond\_broadcast (cond\_t \*c);
  - Wake one/all threads waiting on c

# Condition Vars & Locks

**C/Vs are also used without monitors in conjunction with locks**

**A monitor ≈ a module whose state includes a C/V and a lock**

- Difference is syntactic; with monitors, compiler adds the code

**It is “just as if” each procedure in the module calls acquire() on entry and release() on exit**

- But can be done anywhere in procedure, at finer granularity

**With condition variables, the module methods may wait and signal on independent conditions**

# Condition Vars & Locks

**C/Vs are also used without monitors in conjunction with locks**

**A monitor ≈ a module whose state includes a C/V and a lock**

- Difference is syntactic; with monitors, compiler adds the code

**It is “just as if” each procedure in the module calls acquire() on entry and release() on exit**

- But can be done anywhere in procedure, at finer granularity

**With condition variables, the module methods may wait and signal on independent conditions**

# Condition Vars & Locks

**Why must cond\_wait both release mutex\_t & sleep?**

- void cond\_wait(cond\_t \*c, mutex\_t \*m);

**Why not separate mutexes and condition variables?**

```
while (count == BUFFER_SIZE) {  
    mutex_unlock(&mutex);  
    cond_wait(&not_full);  
    mutex_lock(&mutex);  
}
```

# Condition Vars & Locks

**Why must cond\_wait both release mutex\_t & sleep?**

- void cond\_wait(cond\_t \*c, mutex\_t \*m);

**Why not separate mutexes and condition variables?**

Producer:

```
while (count == BUFFER_SIZE) {  
    mutex_unlock(&mutex);  
  
    cond_wait(&not_full);  
    mutex_lock(&mutex);  
}
```

Consumer:

```
while (count == BUFFER_SIZE) {  
  
    mutex_unlock(&mutex);  
    count --;  
    cond_signal(&not_full);  
    mutex_lock(&mutex);  
}
```

# Monitors and Java

## A lock and condition variable are in every Java object

- No explicit classes for locks or condition variables

## Every object is/has a monitor

- At most one thread can be inside an object's monitor
- A thread enters an object's monitor by
  - Executing a method declared “**synchronized**”
  - Executing the body of a “**synchronized**” statement
- The compiler generates code to acquire the object's lock at the start of the method and release it just before returning
  - The lock itself is implicit, programmers do not worry about it

# Condition Vars & Locks

**Every object can be treated as a condition variable**

- Half of Object's methods are for synchronization!

**Take a look at the Java Object class:**

- Object.wait(\*) is Condition::wait()
- Object.notify() is Condition::signal()
- Object.notifyAll() is Condition::broadcast()

# Next time...

**Read Chapter 15, 16, 18**