

# CE 440 Introduction to Operating System

Lecture 13: Pintos Virtual Memory Notes  
Fall 2025

Prof. Yigong Hu

BOSTON  
UNIVERSITY

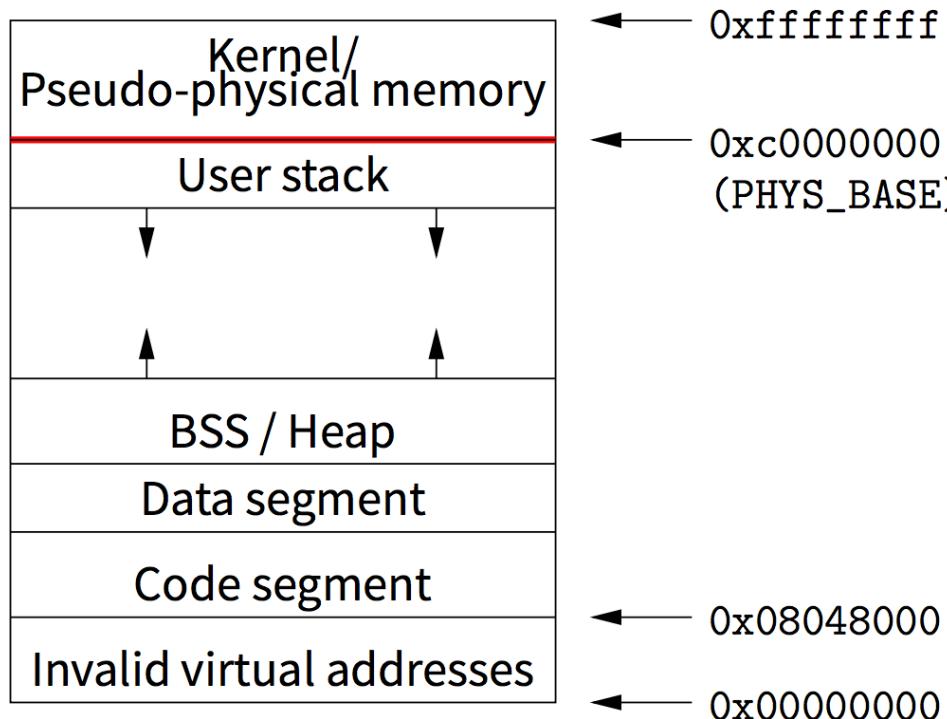
Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

# **Administrivia**

**Lab 2 overview session this Friday**

- 2:30 - 4:00 PM, PHO305

# Pintos Virtual Memory Layout



i.e., the kernel lives in every process' address space!

## A process' virtual address space is split into two regions

- The kernel lives in the high memory region, typically highest 1GB, i.e., from 3 to 4 GB.
- The user memory lives in the lower region, typically lower 3 GB, i.e., from 0 to 3 GB.

# User Virtual Memory

**Per process: a new page directory (pagedir) for each process**

```
struct thread
{
    tid_t tid; /* Thread identifier. */
    enum thread_status status; /* Thread state. */
    char name[16]; /* Name (for debugging purposes). */
    uint8_t *stack; /* Saved stack pointer. */
    int priority; /* Priority. */
    struct list_elem allelem; /* List element for all threads list. */
    struct list_elem elem; /* List element. */

#define USERPROG
    /* Owned by userprog/process.c. */
    uint32_t *pagedir;
#endif /* Page directory. */

/* Owned by thread.c. */
unsigned magic; /* Detects stack overflow. */
};
```

# How Is A User Process Started?

```
$ pintos -p ../../examples/echo -a echo -- -f -q run 'echo ec440'
```

```
static void
run_task (char **argv)
{
    const char *task = argv[1];
    printf ("Executing '%s':\n", task);
#ifndef USERPROG
    process_wait (process_execute (task));
#else
#endif
    run_test (task);
    printf ("Execution of '%s' complete.\n", task);
}
```

# How Is A User Process Started?

```
$ pintos -p ../../examples/echo -a echo -- -f -q run 'echo ec440'
```

```
tid_t  
process_execute (const char *file_name)  
{  
    char *fn_copy;  
    tid_t tid;  
  
    /* Make a copy of FILE_NAME.  
       Otherwise there's a race between the caller and load(). */  
    fn_copy = palloc_get_page (0);  
    if (fn_copy == NULL)  
        return TID_ERROR;  
    strlcpy (fn_copy, file_name, PGSIZE);  
  
    /* Create a new thread to execute FILE_NAME. */  
    tid = thread_create (file_name, PRI_DEFAULT, start_process, fn_copy);  
    if (tid == TID_ERROR)  
        palloc_free_page (fn_copy);  
    return tid;  
}
```

“echo ec440”

Why?

The caller might free the file\_name after this function returns!

,e.g., after you implement `exec`.

# How Is A User Process Started?

```
$ pintos -p ../../examples/echo -a echo -- -f -q run 'echo ec440'
```

```
static void
start_process (void *file_name_)
{
    char *file_name = file_name_;
    struct intr_frame if_;
    bool success;

    /* Initialize interrupt frame and load executable. */
    memset (&if_, 0, sizeof if_);
    if_.gs = if_.fs = if_.es = if_.ds = if_.ss = SEL_UDSEG;
    if_.cs = SEL_UCSEG;
    if_.eflags = FLAG_IF | FLAG_MBS;
    success = load (file_name, &if_.eip, &if_.esp);

    /* If load failed, quit. */
    palloc_free_page (file_name);
    if (!success)
        thread_exit ();

    /* Start the user process by simulating a return from an interrupt */
    asm volatile ("movl %0, %%esp; jmp intr_exit" : : "g" (&if_) : "memory");
    NOT_REACHED ();
}
```

# How Is A User Process Started?

```
$ pintos -p ../../examples/echo -a echo -- -f -q run 'echo ec440'
```

```
bool load (const char *file_name, void (**eip) (void), void **esp)
{
    struct thread *t = thread_current ();
    ...

    /* Allocate and activate page directory. */
    t->pagedir = pagedir_create ();
    if (t->pagedir == NULL)
        goto done;
    process_activate ();

    /* Open executable file. */
    file = filesys_open (file_name);
    ...
}

void process_activate (void)
{
    struct thread *t = thread_current();
    pagedir_activate (t->pagedir);

    /* Set thread's kernel stack for use in processing interrupts. */
    tss_update ();
}
```

```
void pagedir_activate (uint32_t *pd)
{
    if (pd == NULL)
        pd = init_page_dir;
    asm volatile ("movl %0, %%cr3" : : "r" (vtop
(pd)) : "memory");
}
```

After this point, the user virtual  
memory mappings changed!

# Wait

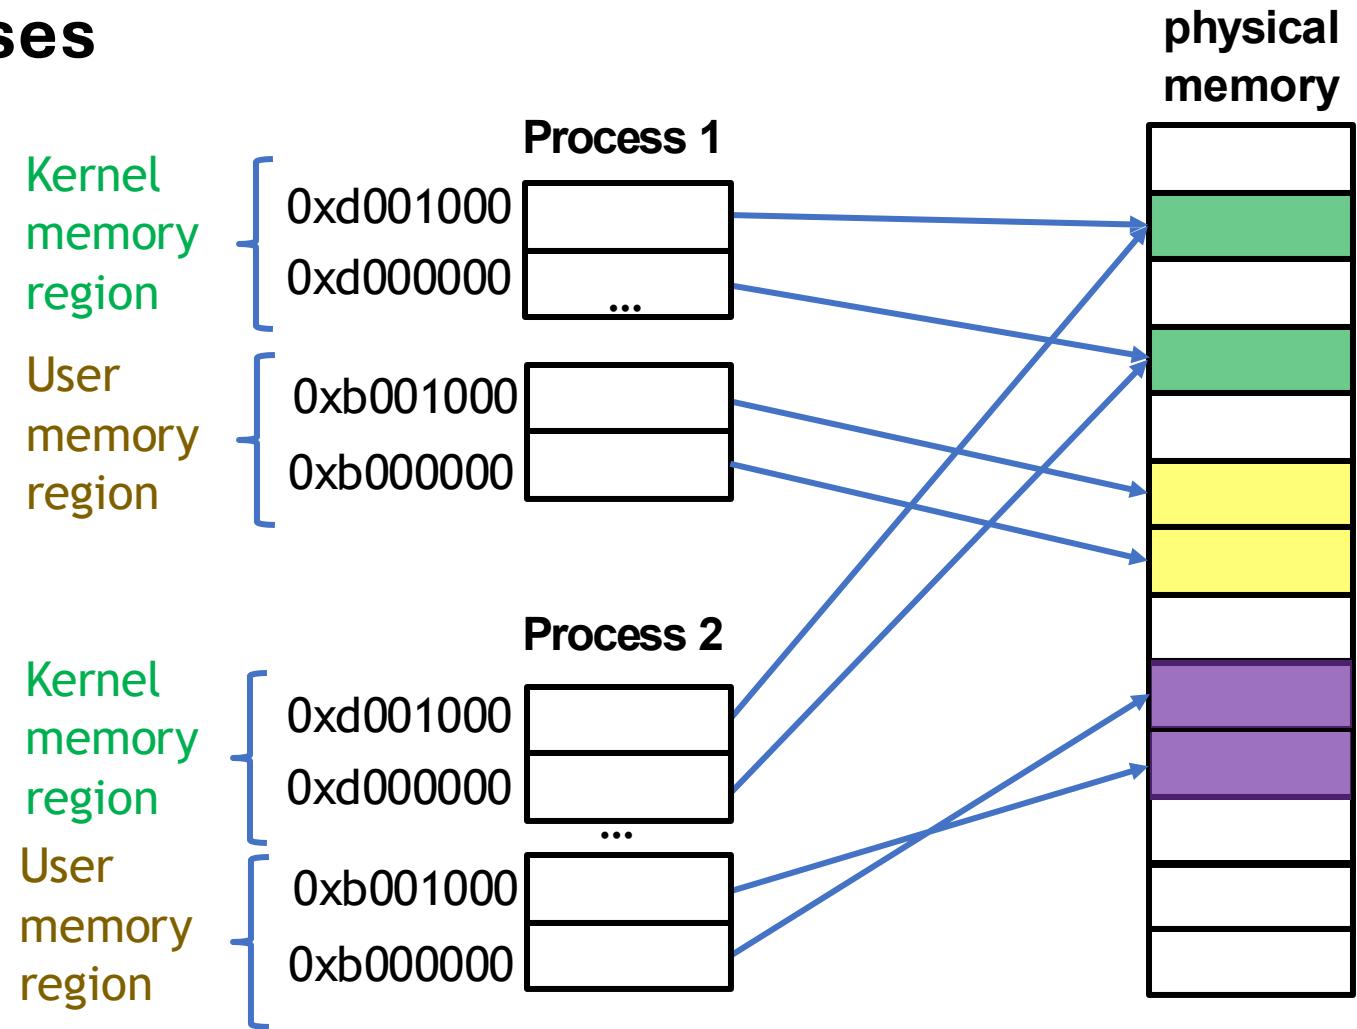
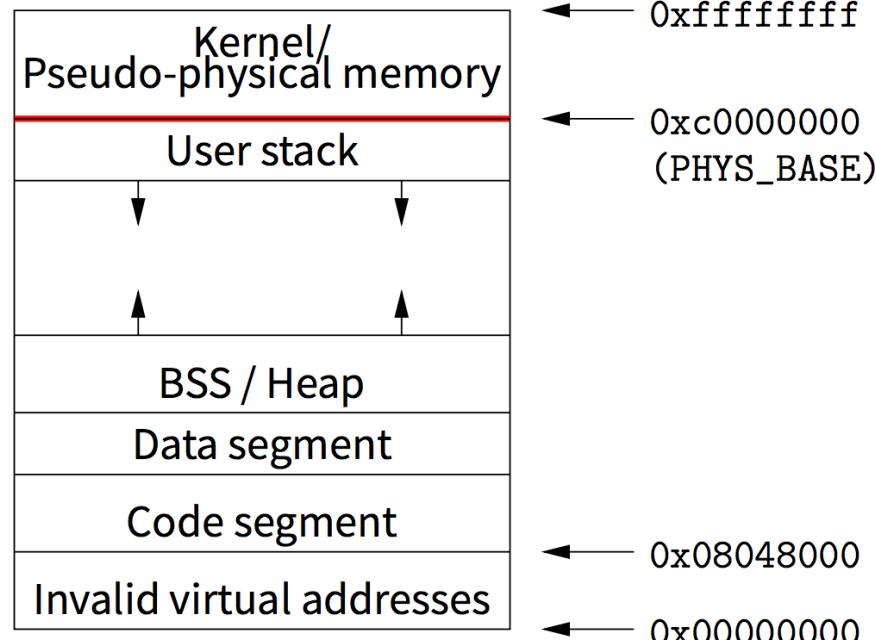
We just changed the user virtual memory mappings, how is it OK for us to still access these variables we created earlier, e.g., file\_name?

A related concern: how to access variables across multiple processes?

- e.g., to implement `int wait (pid_t pid)` you want to create a variable in **struct thread** to store some information for a process,
  - e.g., `thread->wait_status`,
- but how can you read/write this variable from the parent process?

# Answer: We're in the Kernel!

The kernel virtual memory mappings  
are the same across all processes



# Answer: We're in the Kernel!

The kernel virtual memory mappings are the same across all processes

## Implications:

- When we context switch to another process, although it involves changing the page tables, the kernel virtual memory addresses are still valid after the switch
- All objects created in the kernel functions are accessible across processes
  - e.g., **static struct list** all\_list;      threadX->wait\_status
- Memory for user processes will be freed when a user process exits, but memory objects allocated within the kernel code using **malloc** should be explicitly freed!

# How Is This Implemented?

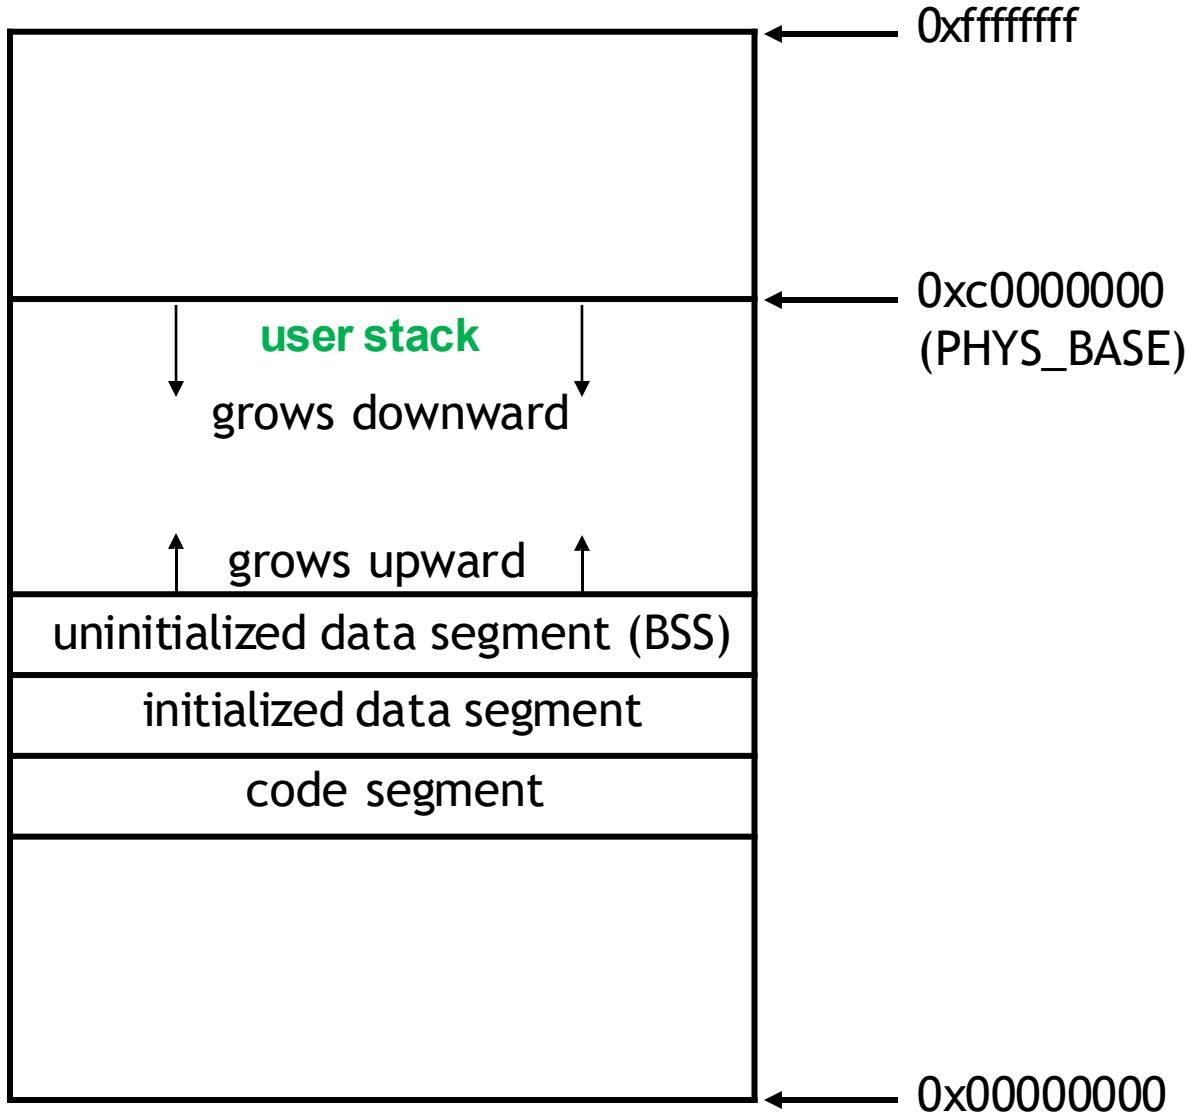
```
bool load (const char *file_name, void (**eip) (void), void **esp)
{
    struct thread *t = thread_current ();
    ...
    /* Allocate and activate page directory. */
    t->pagedir = pagedir_create ();
    if (t->pagedir == NULL)
        goto done;
    process_activate ();
    ...
}

uint32_t *
pagedir_create (void)
{
    uint32_t *pd = palloc_get_page (0);
    if (pd != NULL)
        memcpy (pd, init_page_dir, PGSIZE);
    return pd;
}
```

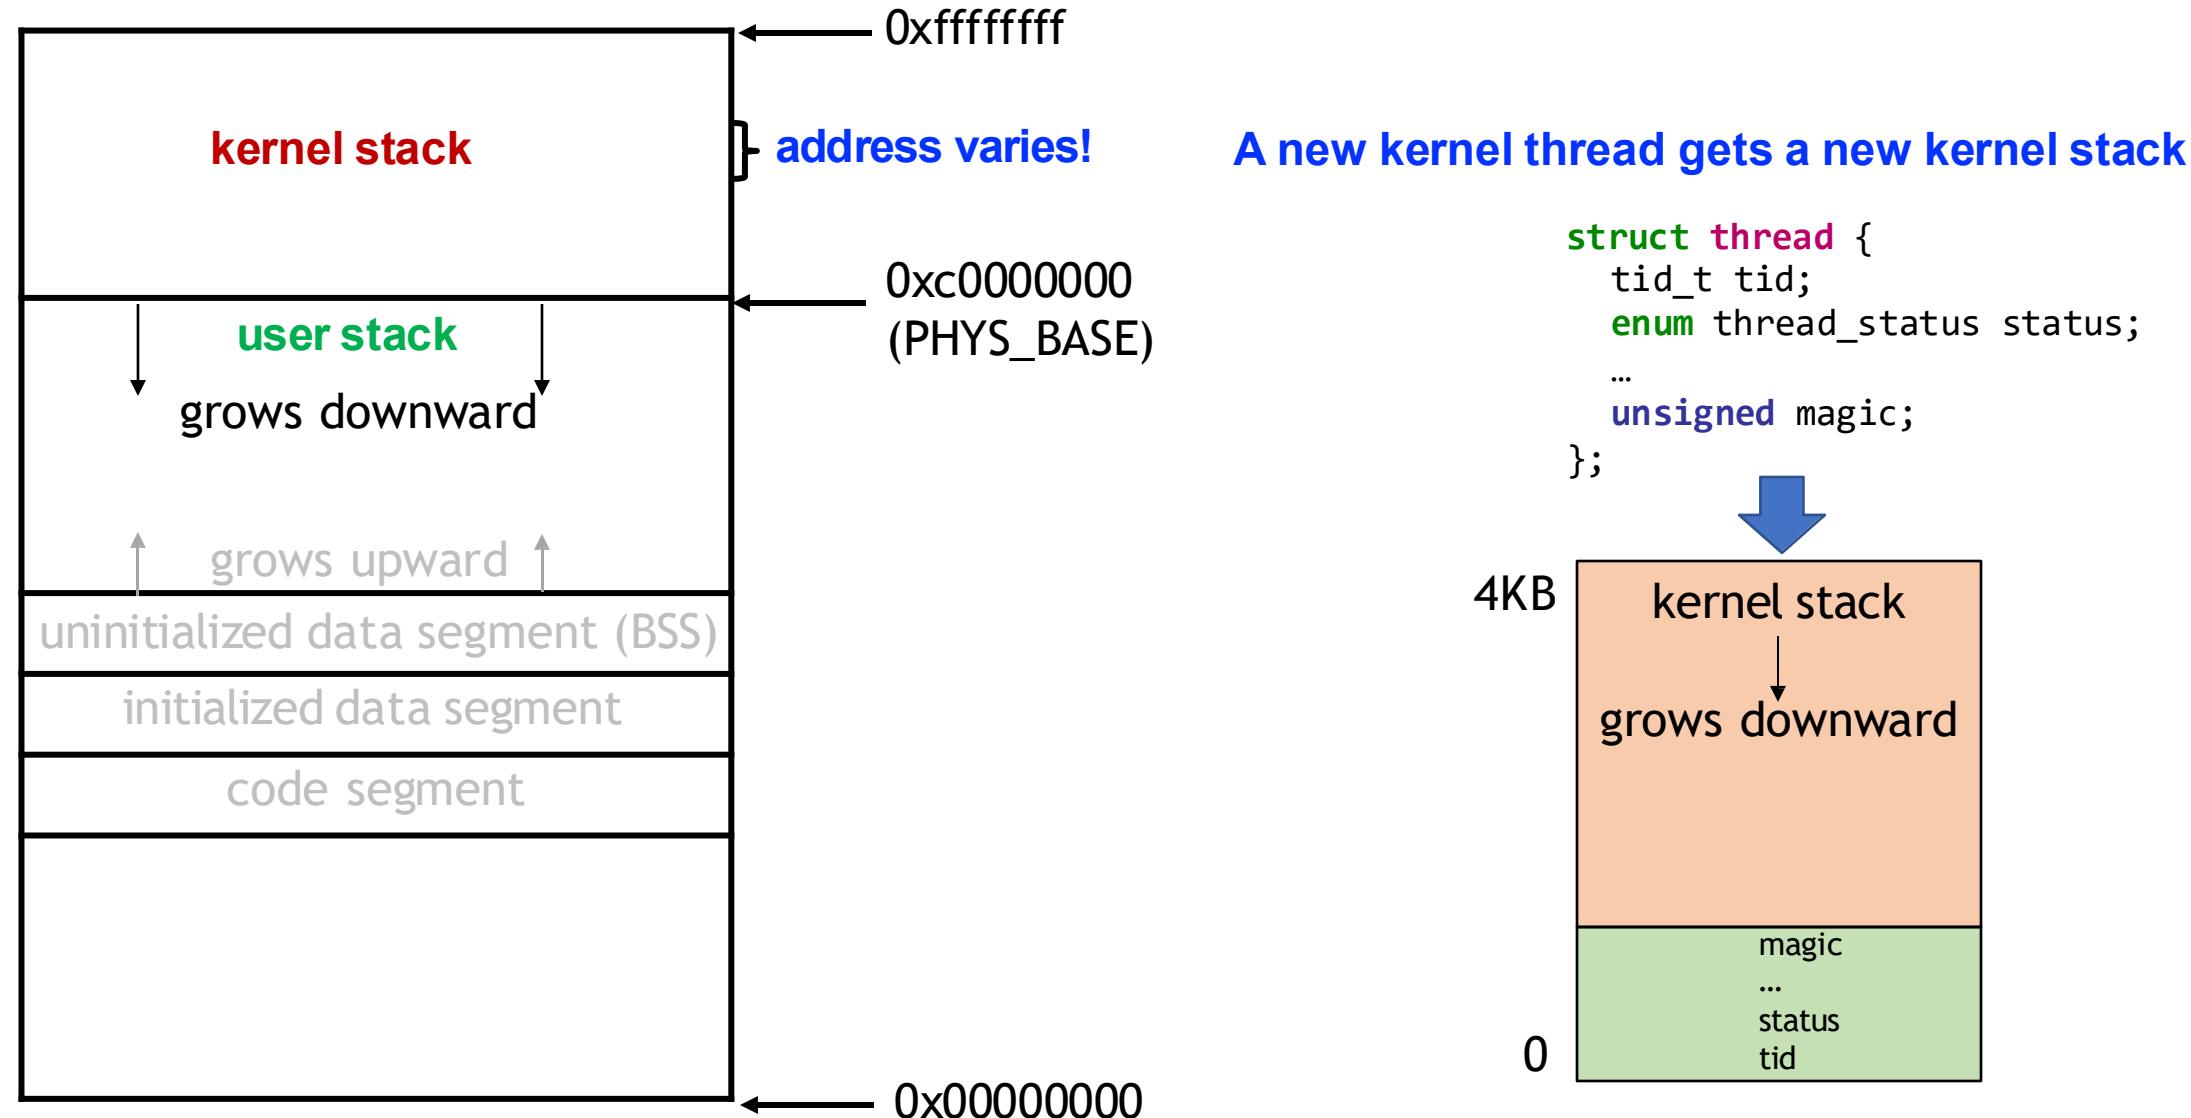
A blue arrow points from the text "Initialized in `paging_init()` in `thread.c`" to the opening brace of the `pagedir_create` function.

Initialized in `paging_init()` in `thread.c`

# User Stack vs Kernel Stack



# User Stack vs Kernel Stack



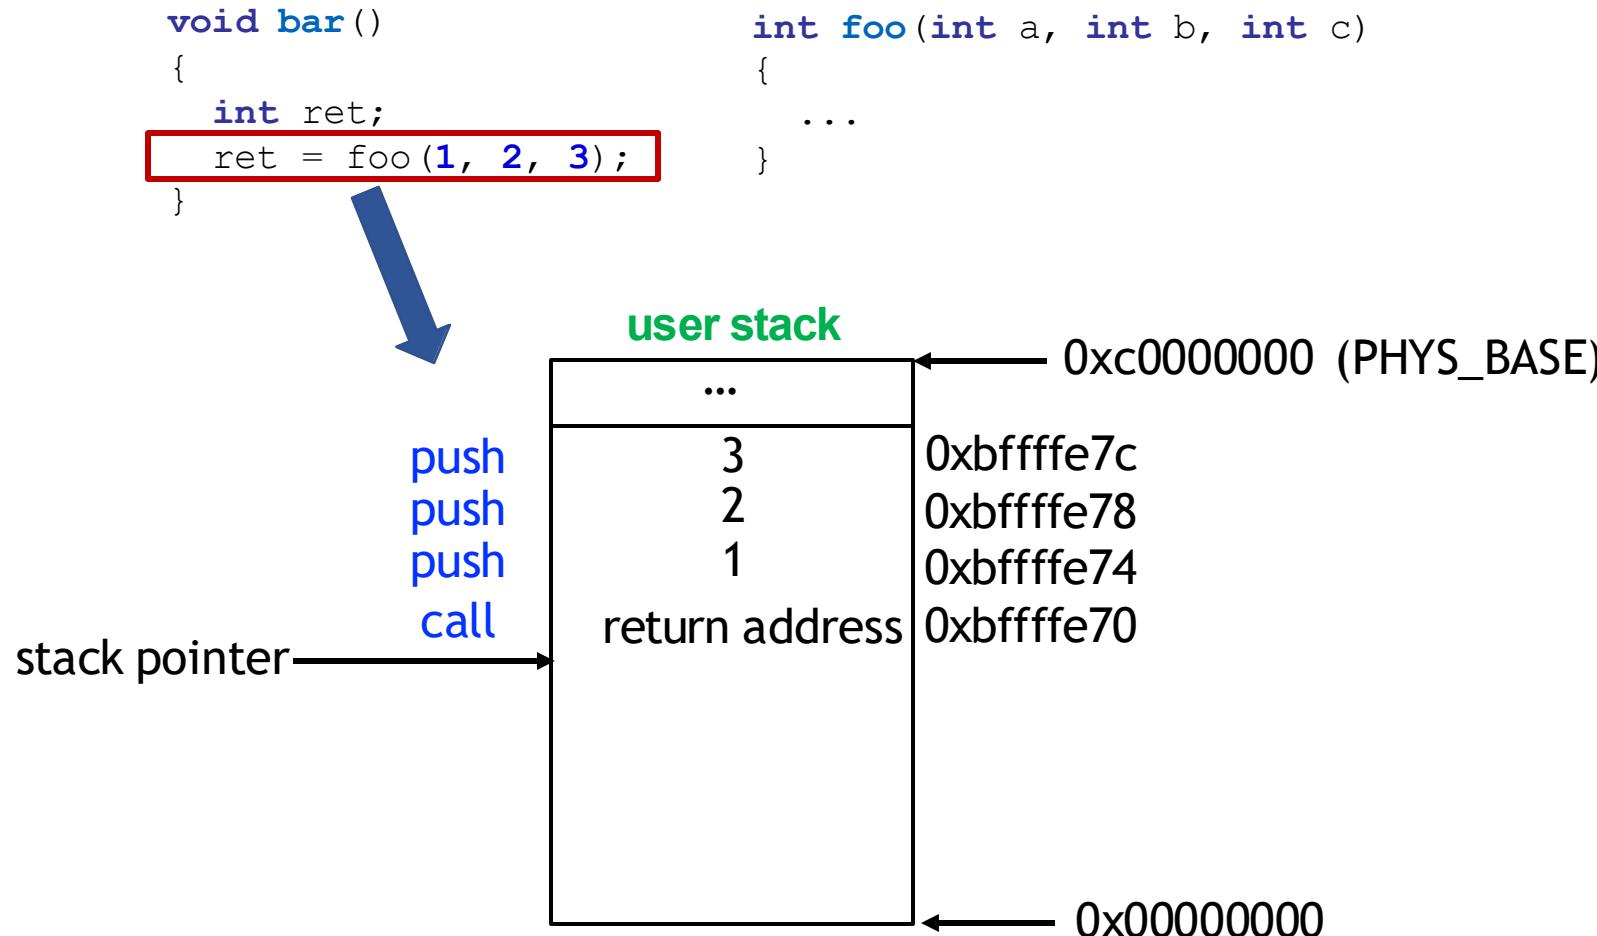
# Lab 2

## Minimal changes to get started:

1. setup\_stack()      \*esp = PHYS\_BASE;    →    \*esp = PHYS\_BASE - 12;
2. change process\_wait() to an infinite loop

# Why Setting esp to PHYS\_BASE – 12 ?

A temporary setup for obeying x86 calling convention



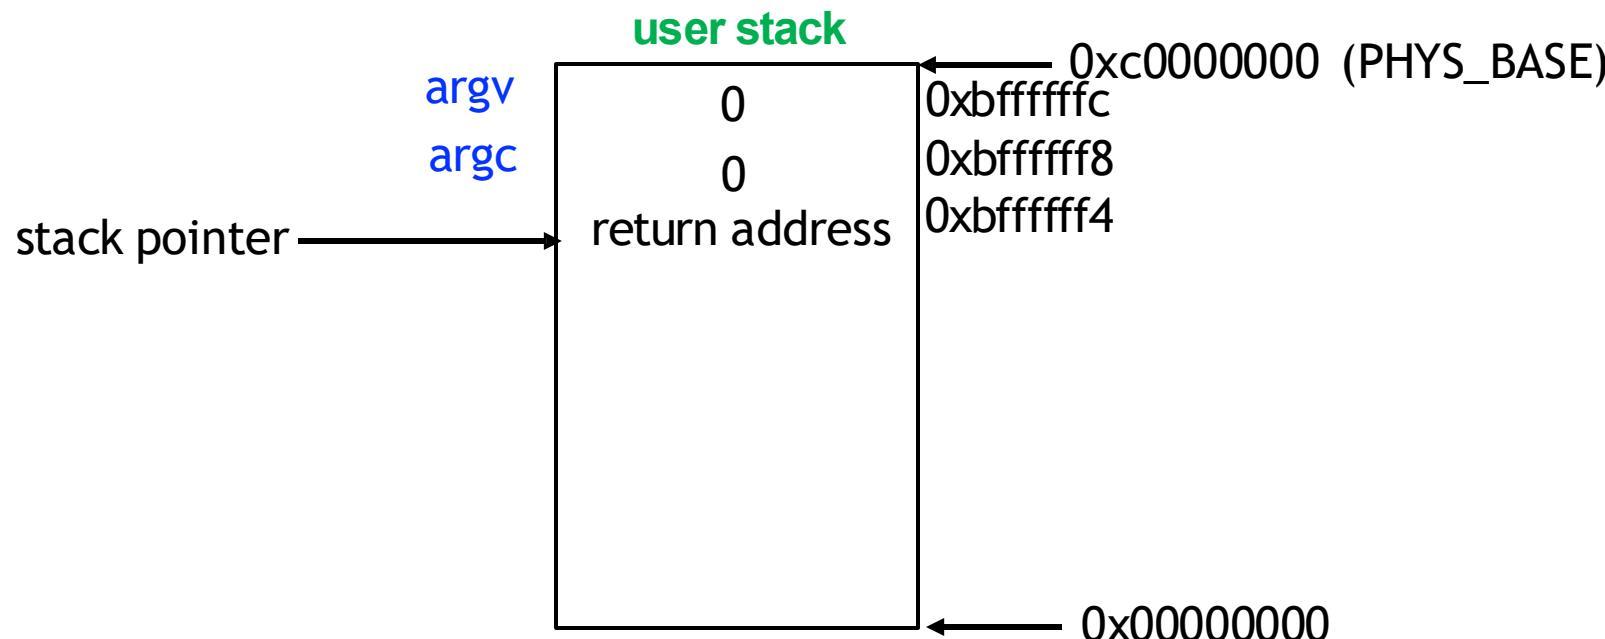
# Why setting esp to PHYS\_BASE - 12 ?

## A temporary setup for obeying x86 calling convention

- Every user program's entry point is:

```
void _start (int argc, char *argv[])
{
    exit (main (argc, argv));
}
```

- Minimal 3 elements on user stack, each 4 bytes = 12



# Why setting esp to PHYS\_BASE - 12 ?

## A temporary setup for obeying x86 calling convention

- Every user program's entry point is:

```
void __start (int argc, char *argv[])
{
    exit (main (argc, argv));
}
```

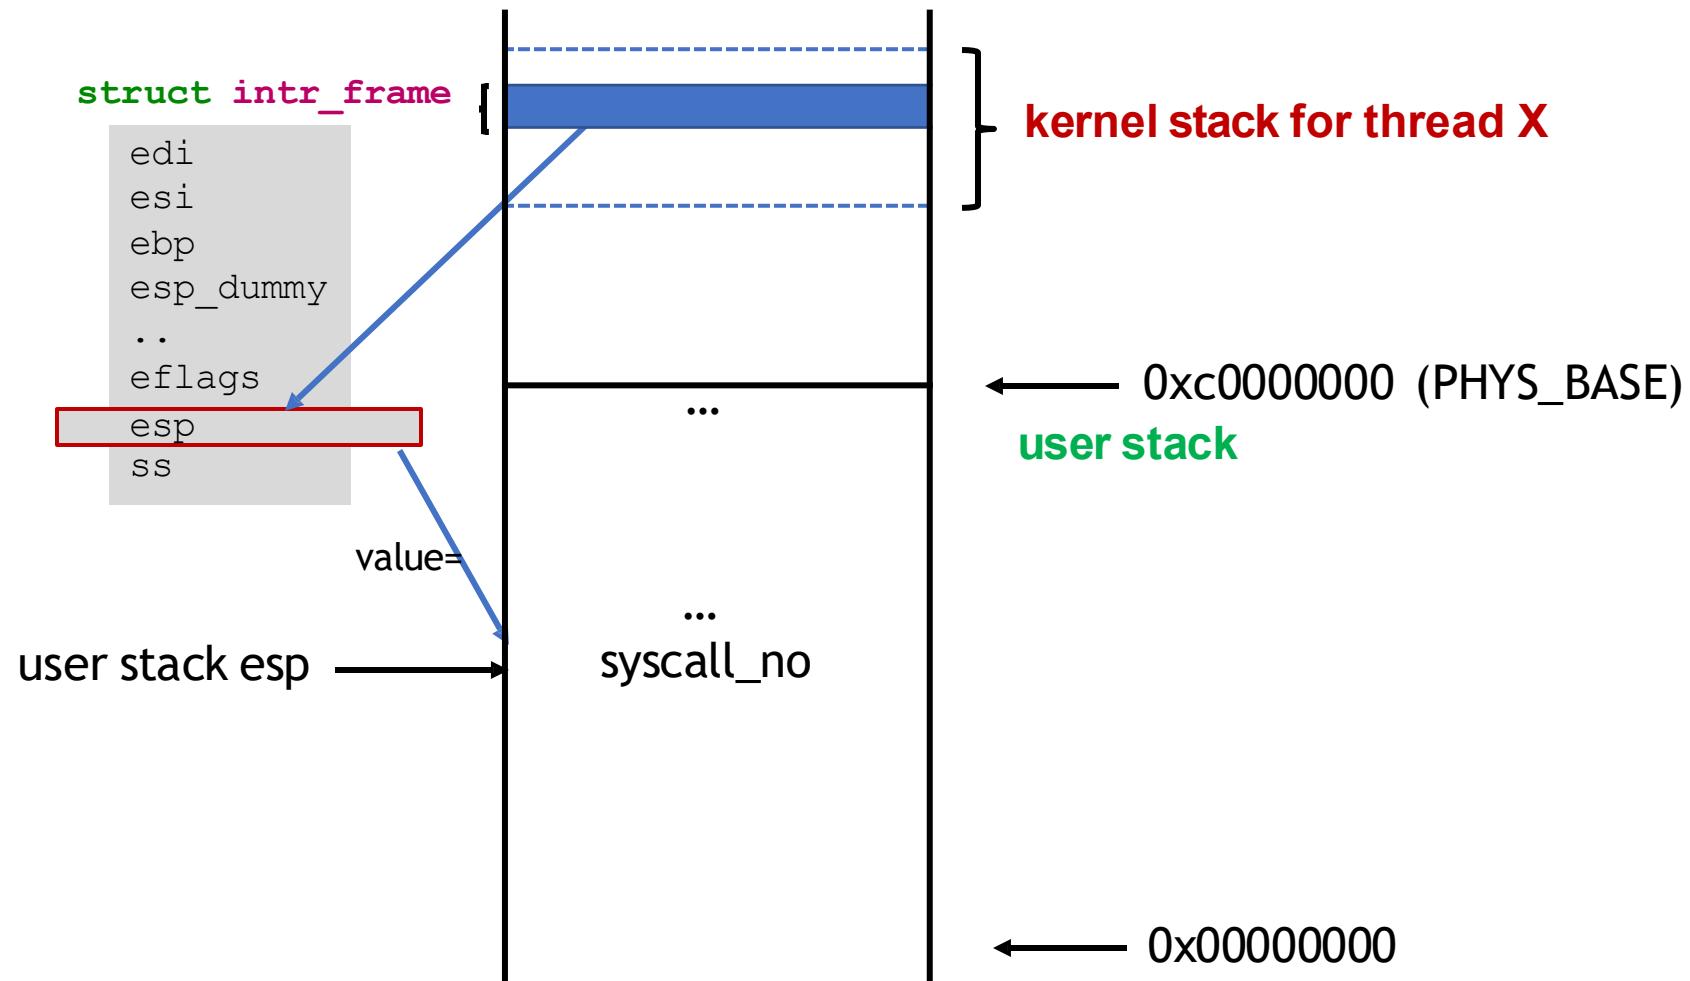
- Minimal 3 elements on user stack, each 4 bytes = 12

### Note: this is only a temporary setup

- Once you implement argument parsing, you should set esp correctly based on the actual arguments pushed on the user stack

# System Call

Through trap (an interrupt frame)



How to retrieve the syscall no  
in `syscall_handler`?

from reading user memory  
at `intr_frame->esp`

# User Memory Access

## Upon system call, **no** page directory switch

- i.e., in syscall\_handler, the kernel can directly access user memory by dereferencing it
- However, must carefully check each user memory address for robustness!

## Two approaches for checking + accessing user memory

- Software approach: using pagedir methods to check validity of an address
  - Easier (straightforward), but slower
- Hardware approach: leveraging page fault to detect invalid address
  - Fast, a bit more difficult to understand (but not difficult to implement once you figure it out)

# Hardware Approach

Try loading the memory from a given address `addr`

- Assume `addr` is the function argument

```
movl 4(%esp), %edx;      ==> edx = addr
movzbl (%edx), %eax;    ==> eax = [addr]
```

- Problem:** we'll get a page fault if `addr` is invalid
- Idea:** let page fault handler inform us, how?

# Hardware Approach

**Use the given helper function, modify page fault handler**

```
/* Reads a byte at user virtual address UADDR.  
UADDR must be below PHYS_BASE.  
Returns the byte value if successful, -1 if a segfault  
occurred. */  
  
static int get_user (const uint8_t *uaddr)  
{  
    int result;  
    asm ("movl $1f, %0; movzbl %1, %0; 1:"  
         : "=a" (result) : "m" (*uaddr));  
    return result;  
}
```



```
get_user:  
    movl 4(%esp), %edx;  
    movl $1f, %eax;  
    movzbl (%edx), %eax;  
1:  
    ret
```

- If addr is valid, eax has the value
- If addr is invalid, the page fault handler will
  - set eip to address of label 1 (stored in eax now)
  - set eax to be -1 (0xffffffff);
  - resume to ret

# Hardware Approach

**Use the given helper function, modify page fault handler**

```
/* Reads a byte at user virtual address UADDR.  
 UADDR must be below PHYS_BASE.  
 Returns the byte value if successful, -1      if a segfault  
 occurred. */  
static int get_user (const uint8_t *uaddr)  
{  
    int result;  
    asm ("movl $1f, %0; movzbl %1, %0; 1:"  
        : "=a" (result) : "m" (*uaddr));  
    return result;  
}
```

**But what if the value at uaddr is -1? We can't tell if it's invalid or not!**

- **Solution:** read one byte at a time!
  - If value is valid, at most can be 255 (0xff)
  - How to represent a valid -1? Read four bytes (call `get_user` four times), convert to an integer!