

# CE 440 Introduction to Operating System

## Lecture 16: File System Fall 2025

Prof. Yigong Hu



Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

# Administrivia

**Lab 3a and 3b is out**

- Start the project early

**Homework 4 is out**

# File System ~~Not~~ Fun

## File systems: a challenging OS design topic

- More papers on FSes than any other single topic

## Main tasks of file system:

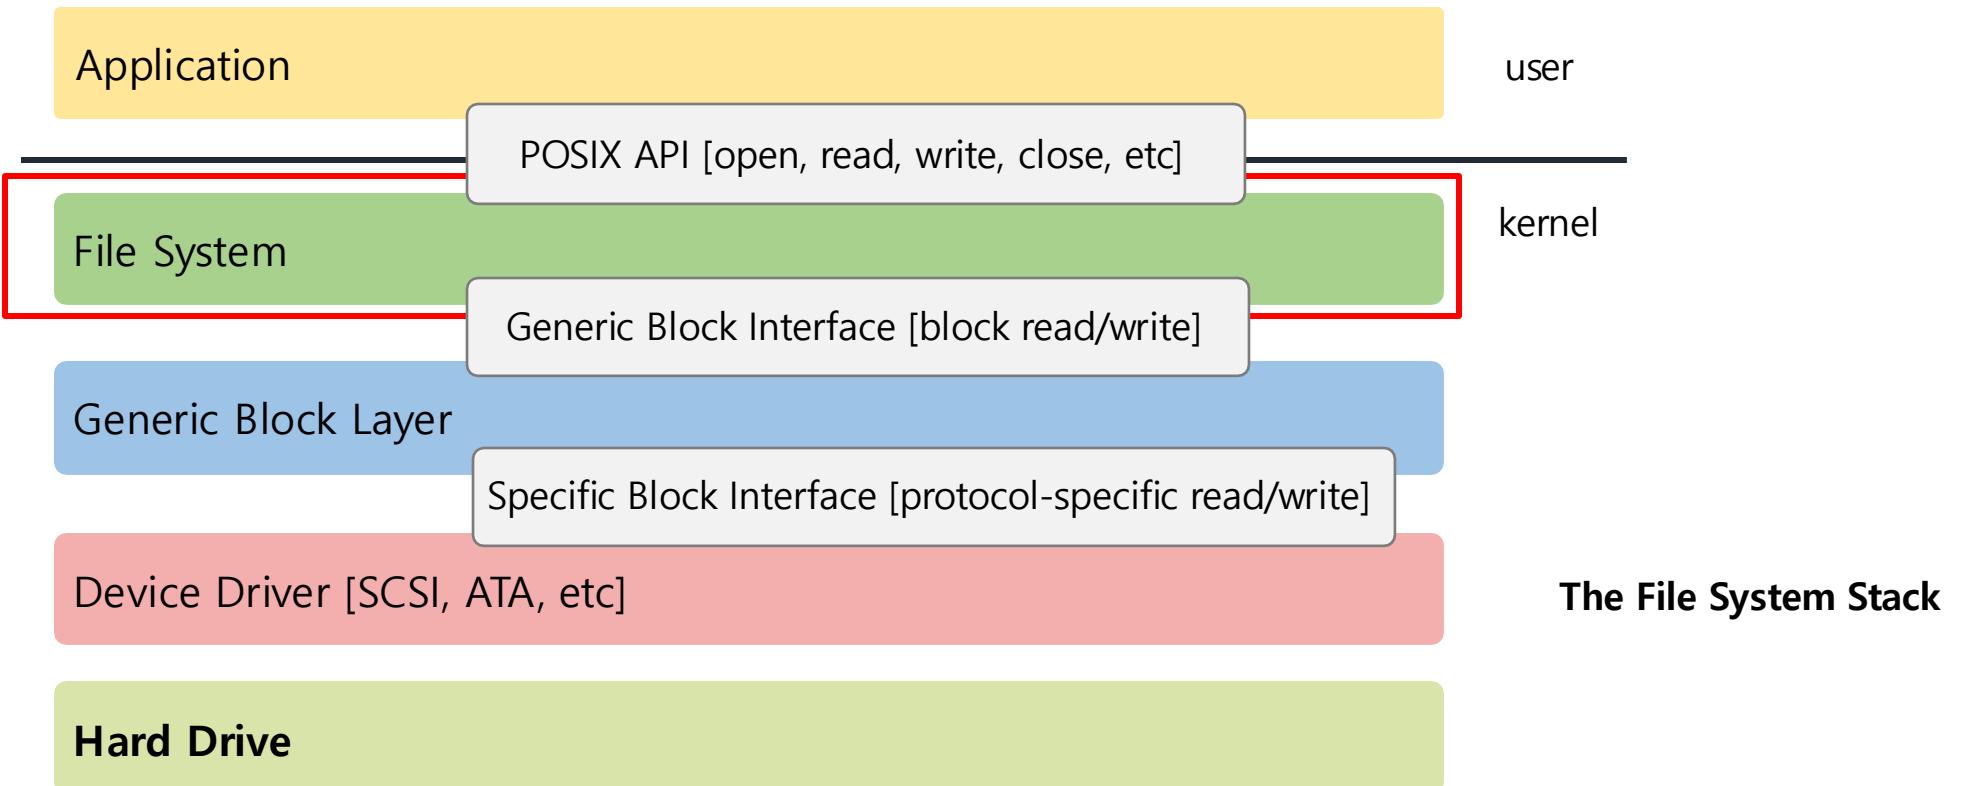
- Don't go away (ever)
- Associate bytes with name (files)
- Associate names with each other (directories)
- Can implement file systems on disk, over network, in memory, in non-volatile ram (NVRAM), on tape, w/ paper.
- We'll focus on disk and generalize later

## Today: files, directories

# File System Abstraction

File system **specifics** of which disk class it is using.

- It issues **block read** and **write** request to the generic block layer.



# Files

## File: named bytes on disk

- Data with some properties
- Contents, size, owner, last read/write time, protection, etc.

## How is a file's data managed by the file system?

- Next lecture's topic
- Basic idea (in Unix): a struct called an index node or inode
  - **Describe where on the disk the blocks for a file are placed**
  - Disk stores an array of inodes, inode # is the index in this array

# File Types

## A file can also have a type

- Understood by the file system
  - Block, character, device, portal, link, etc.
- Understood by other parts of the OS or runtime libraries
  - Executable, dll, source, object, text, etc.

## A file's type can be encoded in its name or contents

- Windows encodes type in name (.com, .exe, .bat, .dll, .jpg, etc.)
- Unix encodes type in contents (magic numbers, initial characters, e.g., #! for shell scripts)

# Basic File Operations

## Unix

- `creat(name)`
- `open(name, how)`
- `read(fd, buf, len)`
- `write(fd, buf, len)`
- `sync(fd)`
- `seek(fd, pos)`
- `close(fd)`
- `unlink(name)`

## Windows

- `CreateFile(name, CREATE)`
- `CreateFile(name, OPEN)`
- `ReadFile(handle, ...)`
- `WriteFile(handle, ...)`
- `FlushFileBuffers(handle, ...)`
- `SetFilePointer(handle, ...)`
- `CloseHandle(handle, ...)`
- `DeleteFile(name)`
- `CopyFile(name)`
- `MoveFile(name)`

# File Access Methods

FS usually provides different file access methods:

- Sequential access
  - read bytes one at a time, in order
  - by far the most common mode
- Random access
  - random access given block/byte number
- Record access
  - file is array of fixed- or variable-length records
  - read/written sequentially or randomly by record #
- Indexed access
  - file system contains an index to a particular field of each record in a file
  - reads specify a value for that field and the system finds the record via the index

What file access method does Unix, Windows provide?

# Directories

## Problem: referencing files

**Users remember where on disk their files are (disk sector no.)?...**

- E.g., like remembering your social security or bank account #

**...People want human digestible names**

**Directories serve two purposes**

- For users, they provide a structured way to organize files
- For FS, they provide a convenient naming interface that allows the separation of logical file organization from physical file placement on the disk

# A Short History of Directories

## Approach 1: Single directory for entire system

- Put directory at known disk location. **If one user uses a name, no one else can**
- Many ancient personal computers work this way

## Approach 2: Single directory for each user

- Still clumsy, and running `ls` on 10,000 files is a real pain

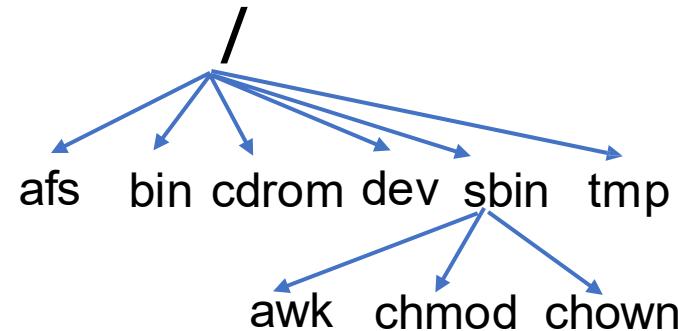
## Approach 3: Hierarchical name spaces

- Allow directory to map names to files or other dirs
- File system forms a tree (or graph, if links allowed)

# Hierarchical Directory

**Used since CTSS (1960s)**

- Unix picked up and used really nicely



**Large name spaces tend to be hierarchical**

- ip addresses, domain names, scoping in programming languages, etc.

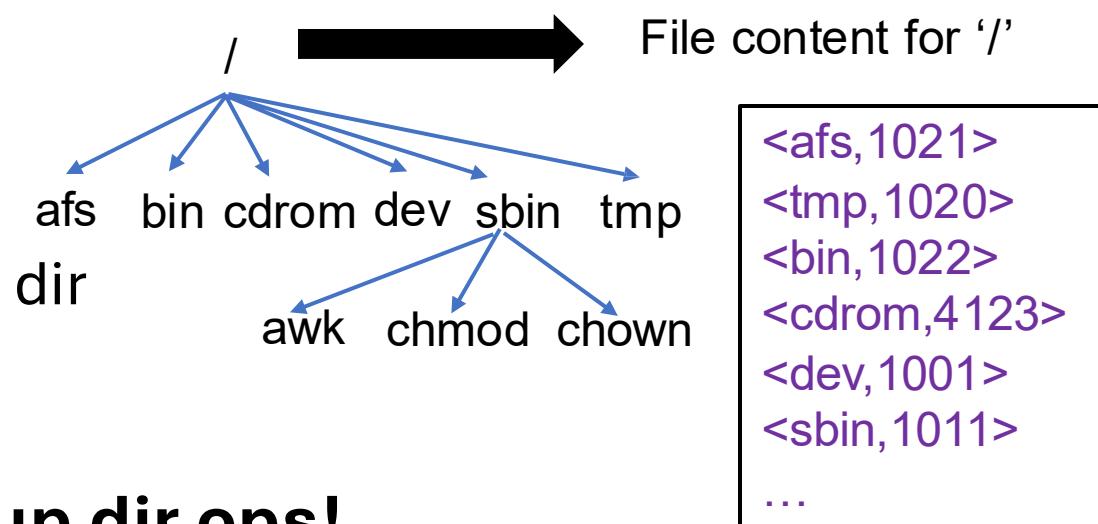
# Directory Internals

A directory is a list of entries

- <name, location> tuple, location is typically the *inode #* (more next lecture)
- An inode describes where on the disk the blocks for a file are placed

Directories stored on disk just like regular files

- File type set to directory
- User's can read just like any other file
- Only special syscalls can write (why?)
- File pointed to by the location may be another dir
- Makes FS into hierarchical tree



Simple, plus speeding up file ops speeds up dir ops!

# Path Name Translation

Let's say you want to open “/one/two/three.txt”

What does the file system do?

- Directory entries map file names to location ([inode #](#))
- Open directory “/”: Where? **Root directory is always inode #2**
- Search for the entry “one”, get location of “one” (in dir entry)
- Open directory “one”, search for “two”, get location of “two”
- Open directory “two”, search for “three”, get location of “three”
- Open file “three”

# Naming Magic

## Bootstrapping: Where do you start looking?

- Root directory always inode #2 (0 and 1 historically reserved)

## Special names:

- Root directory: “/”
- Current directory: “.”
- Parent directory: “..”

## Some special names are provided by shell, not FS:

- User’s home directory: “~”
- Globbing: “foo.\*” expands to all files starting “foo.”

## Using the given names, only need two operations to navigate the entire name space:

- cd name: move into (change context to) directory name
- ls: enumerate all names in current directory (context)

# Basic Directory Operations

## Unix

### Directories implemented in files

- Use file ops to create dirs.

### C library provides a higher-level abstraction for reading directories

- `opendir(name)`
- `readdir(DIR)`
- `seekdir(DIR)`
- `closedir(DIR)`

## Windows

### Explicit directory operations

- `CreateDirectory(name)`
- `RemoveDirectory(name)`

### Very different method for reading directory entries

- `FindFirstFile(pattern)`
- `FindNextFile()`

# **Default Context: Working Directory**

## **Cumbersome to constantly specify full path names**

- In Unix, each process has a “current working directory” (cwd)
- File names *not* beginning with “/” are assumed to be relative to cwd; otherwise translation happens as before

## **Shells track a default list of active contexts**

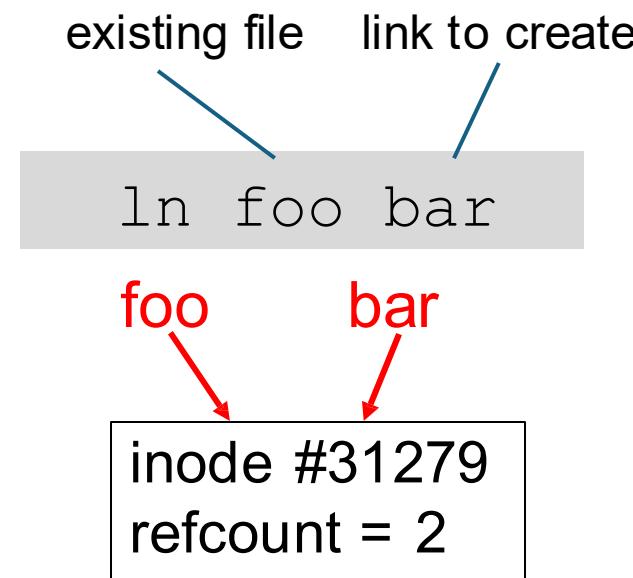
- A “search path” for programs you run
- Given a search path **A:B:C**, the shell will check in A, then B, then C
- Can escape using explicit paths: “./foo”

## **Example of locality**

# Hard Links

**More than one dir entry can refer to a given file**

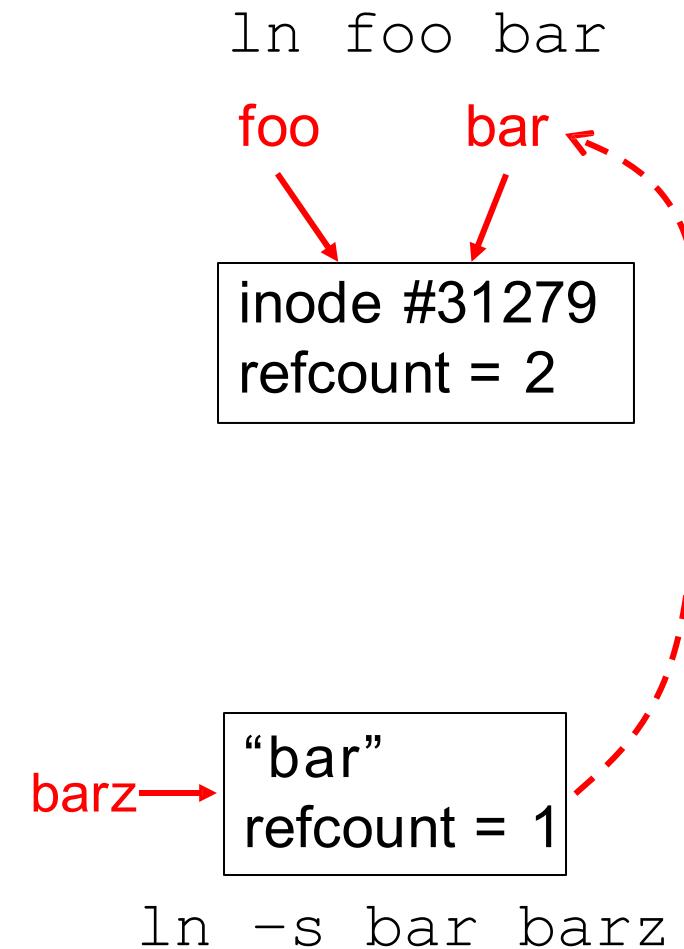
- Hard link creates a synonym for file
- Unix stores count of pointers (“hard links”) to inode
- If one of the links is removed (e.g., rm), the data are still accessible through any other link that remains
- If all links are removed, the space occupied by the data is freed.



# Soft Links

## Soft/symbolic links = synonyms for names

- Point to a file/dir name, but object can be deleted from underneath it (or never exist).
- Unix implements like directories: inode has special “symlink” bit set and contains name of link target
- When the file system encounters a soft link it automatically translates it (if possible).



# File Sharing

## File sharing has been around since timesharing

- Easy to do on a single machine
- PCs, workstations, and networks get us there (mostly)

## File sharing is important for getting work done

- Basis for communication and synchronization

## Two key issues when sharing files

- Semantics of concurrent access
  - What happens when one process reads while another writes?
  - What happens when two processes open a file for writing?
  - **What are we going to use to coordinate?**
- Protection

# Protection

## File systems implement a protection system

- Who can access a file
- How they can access it

## More generally...

- Objects are “what”, subjects are “who”, actions are “how”

A protection system dictates whether a given **action** performed by a given **subject** on a given **object** should be allowed

- You can read and/or write your files, but others cannot
- You can read “/etc/motd”, but you cannot write it

# Representing Protection

## Access Control Lists (ACL)

For each **object**, maintain a list of **subjects** and their permitted actions

## Capabilities

For each **subject**, maintain a list of **objects** and their permitted actions

	Objects		
	/one	/two	/three
Alice	rw	-	rw
Bob	w	-	r
Charlie	w	r	rw

Subjects

ACL

Capability

# ACLs and Capabilities

Approaches differ only in how the table is represented

**Capabilities are easier to transfer**

- They are like keys, can handoff, does not depend on subject

**In practice, ACLs are easier to manage**

- Object-centric, easy to grant, revoke
- To revoke capabilities, have to keep track of all subjects that have the capability  
– a challenging problem

**ACLs have a problem when objects are heavily shared**

- The ACLs become very large
- Use groups (e.g., Unix)

# Unix File Protection

## What approach does Unix use in the FS?

- Answer: both

## ACL: Unix file permissions

## Capability: file descriptors

## How are they used together?

- Conversion through open() system call

Converted to  
capability

```
int fd = open("file.txt", O_WRONLY);  
if (fd == -1)  
    exit(-1);
```

ACL check, expensive

```
for (int i = 0; i < 100; i++)  
    write(fd, buf + i * 4, 4);
```

Use capability from then on

# Summary

## Files

- Operations, access methods

## Directories

- Operations, using directories to do path searches
- 

## Sharing

## Protection

- ACLs vs. capabilities

# Next Chapter

Read Chapter 40, 41

.