

CE 440 Introduction to Operating System

Lecture 20: File System Crash Consistency
Fall 2025

Prof. Yigong Hu

BOSTON
UNIVERSITY

Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

Administrivia

Lab 3a

- Due today (12/01) 11:59 pm
- If you design lab 3a well, 3b is relatively easy

Midterm 2

- Next Wednesday (12/10) in class
- Same format as midterm 1

More on Midterm 2

Covers material in the second half of the class

- Most questions about lecture 10 to lecture 17
- A few basic questions for lecture 18 to lecture 20

Closed book, two double-sided 8.5”x11” pages of notes

Can use a calculator but no other electronic devices

Based upon lecture (textbook), homework, and project

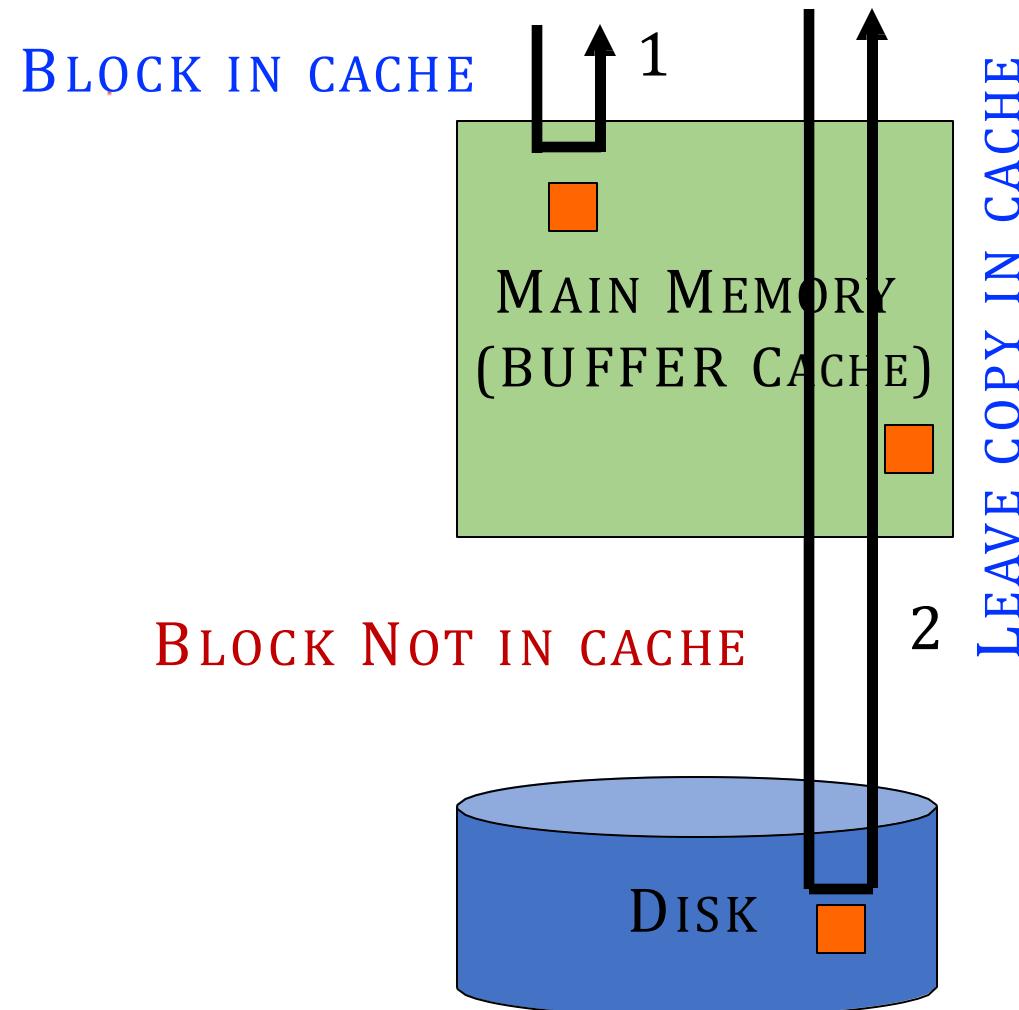
- Do the homework to practice for the exam

Review: File I/O Path (Reads)

File system uses buffer cache to speed up I/O

read() from file

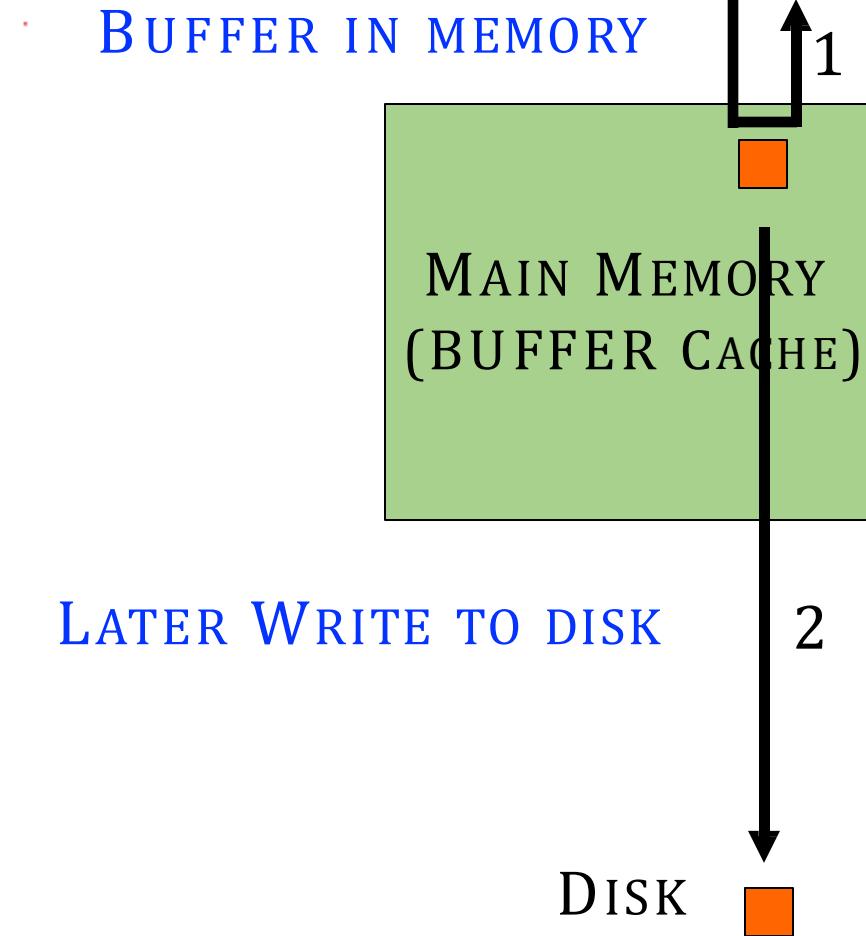
- Check if block is in cache
- If so, return block to user [1 in figure]
- If not, read from disk, insert into cache, return to user [2]



Review: File I/O Path (Writes)

write() to file

- Write is buffered in memory (“write behind”) [1]
- Sometime later, OS decides to write to disk [2]
- Periodic flush or fsync call



Why delay writes?

- Implications for performance
- Implications for reliability

The Consistent Update Problem

Goal:

- Atomically update file system from one **consistent** state to another
- What do we mean by consistent state?

Challenge:

- An update may require modifying several sectors, despite that the **disk only provides atomic write of one sector at a time**

Example: File Creation of /a.txt

Initial state

MEMORY



DISK

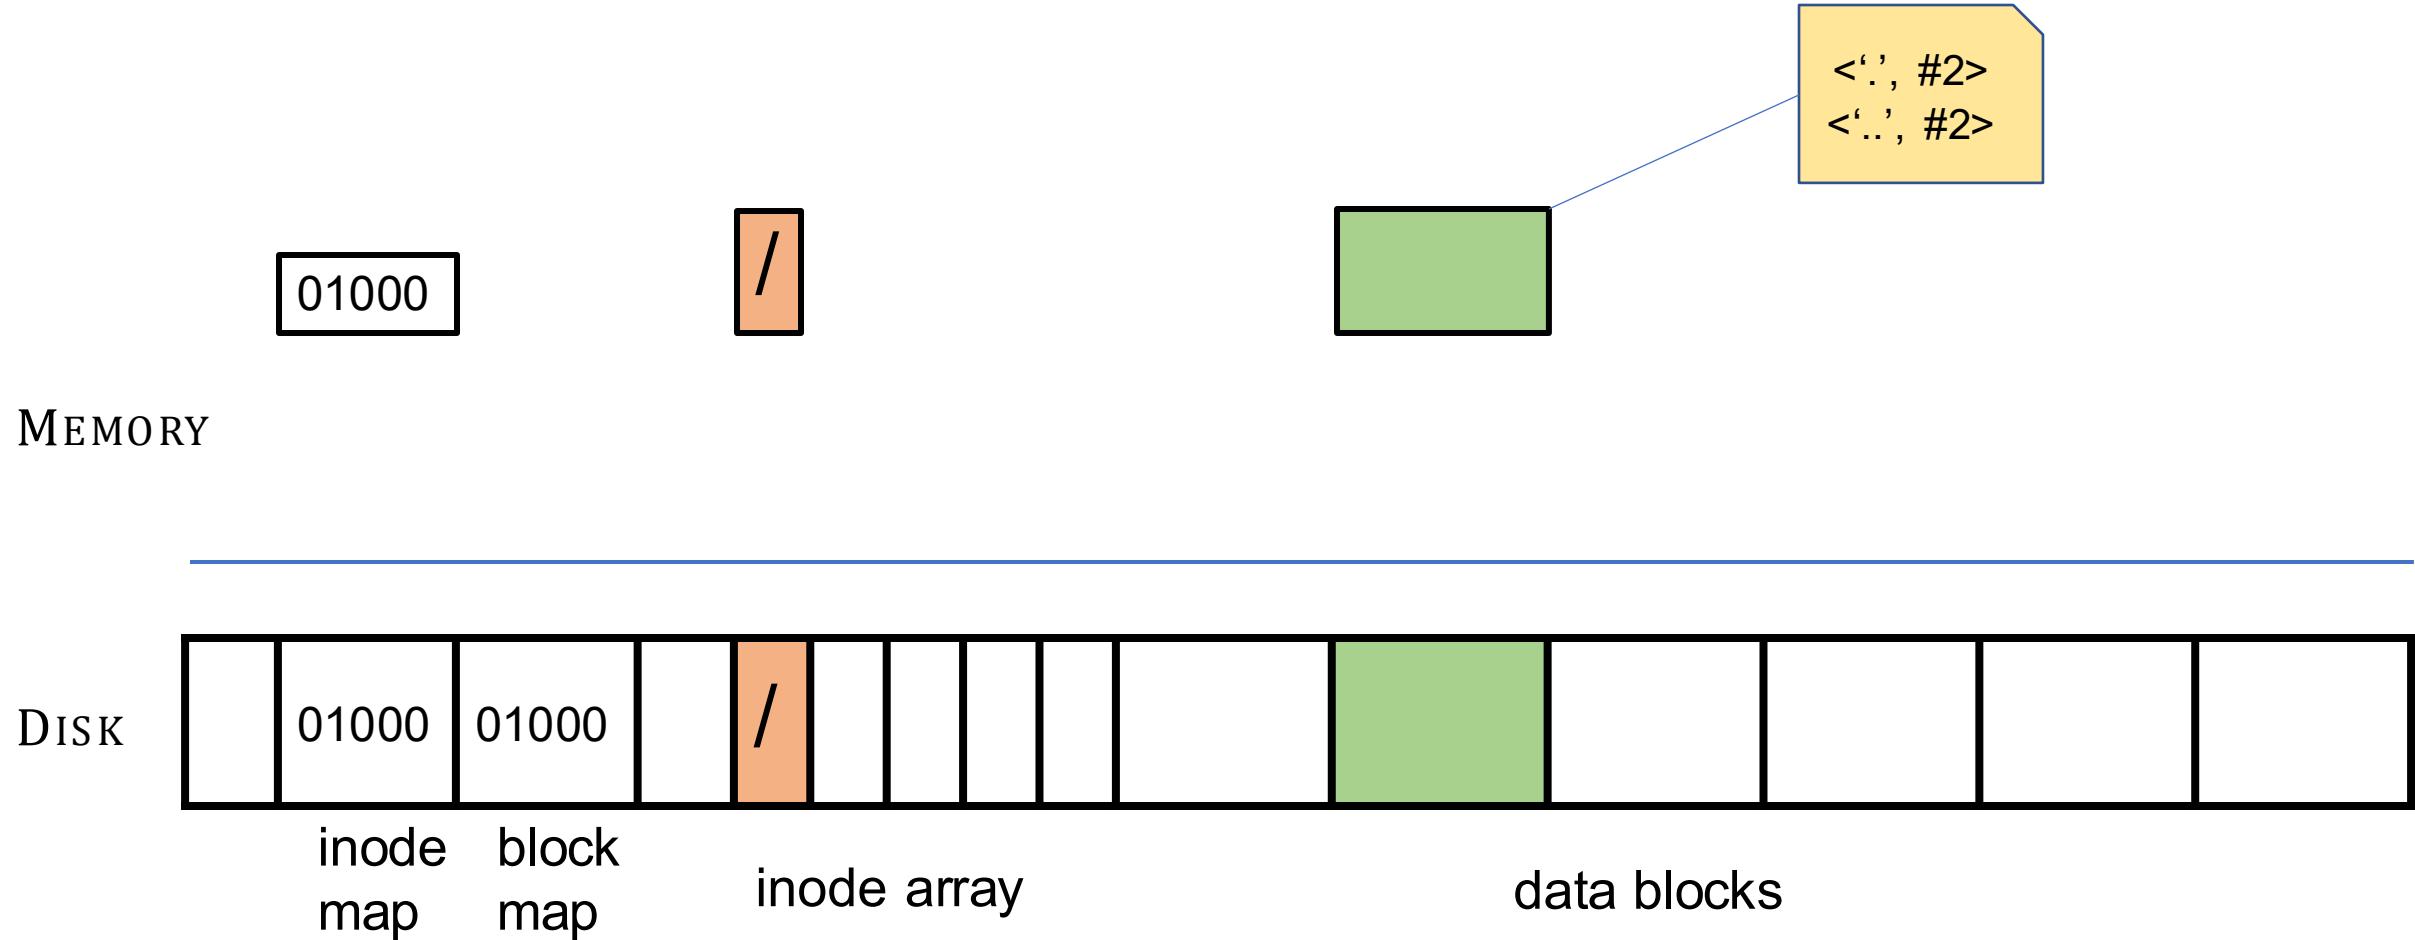
inode
map block
map

inode array

data blocks

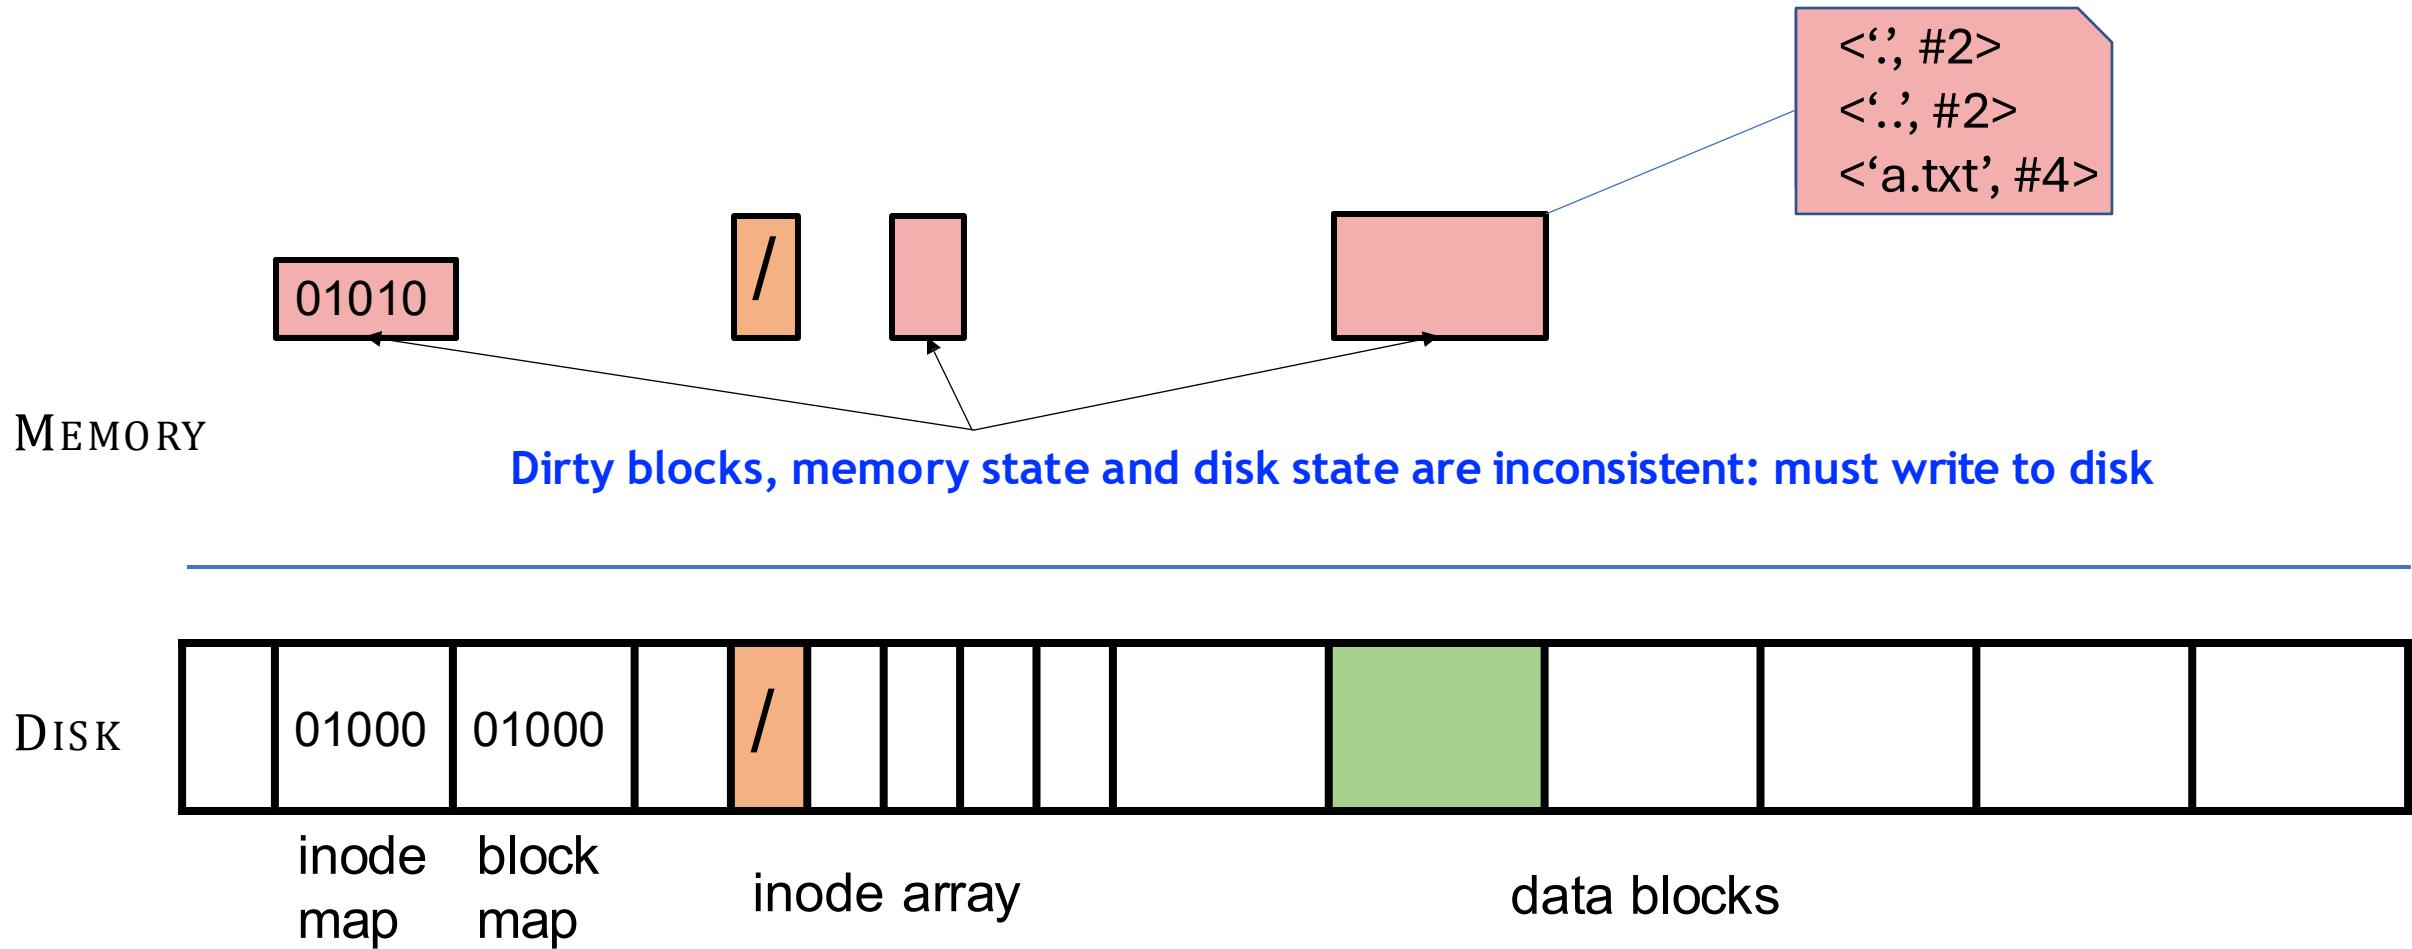
Example: File Creation of /a.txt

Read to in-memory Cache



Example: File Creation of /a.txt

Modify metadata and blocks



Crash?

Disk: atomically write one sector

- Atomic: if crash, a sector is either completely written, or none of this sector is written

An FS operation may modify multiple sectors

Crash → FS partially updated

Possible Crash Scenarios

File creation dirties three blocks

- inode bitmap (B)
- inode for new file (I)
- parent directory data block (D)

Old and new contents of the blocks

- $B = 01000$ $B' = 01010$
- $I = \text{free}$ $I' = \text{allocated, initialized}$
- $D = \{\}$ $D' = \{\langle 'a.txt', 4 \rangle\}$

Possible Crash Scenarios

Crash scenarios: any subset can be written

- B I D
- B' I D
- B I' D
- B I D'
- B' I' D
- B' I D'
- B I' D'
- B' I' D'

The General Problem

Writes: Have to update disk with N writes

- Disk does only a single write atomically

Crashes: System may crash at arbitrary point

- Bad case: In the middle of an update sequence

Desire: To update on-disk structures **atomically**

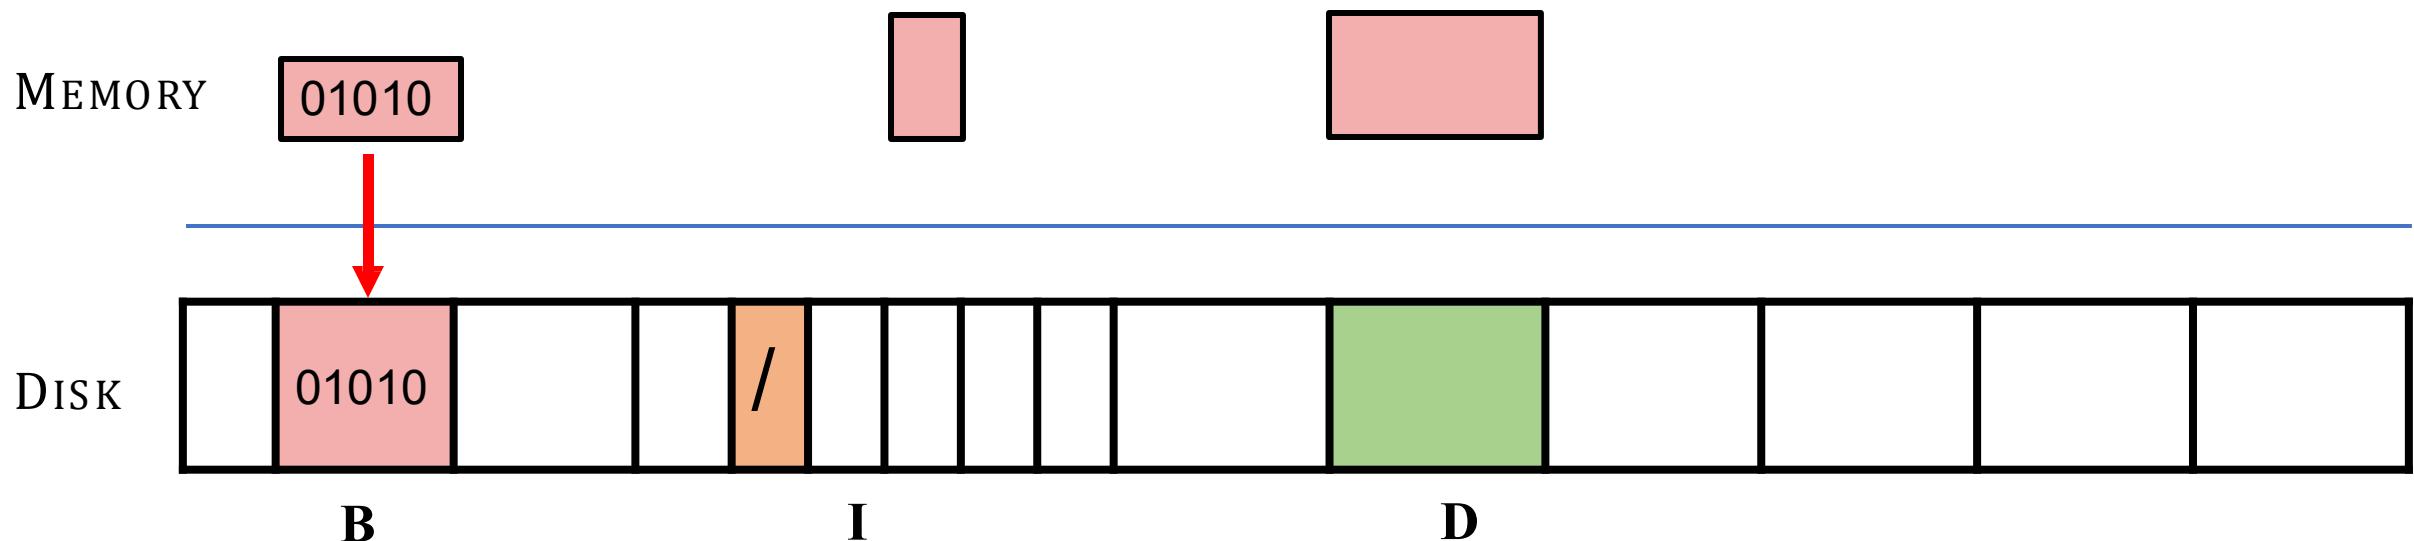
- Either all should happen or none

Example: Bitmap First

Write Ordering: Bitmap (B), Inode (I), Data (D)

- But CRASH after B has reached disk, before I or D

Result?

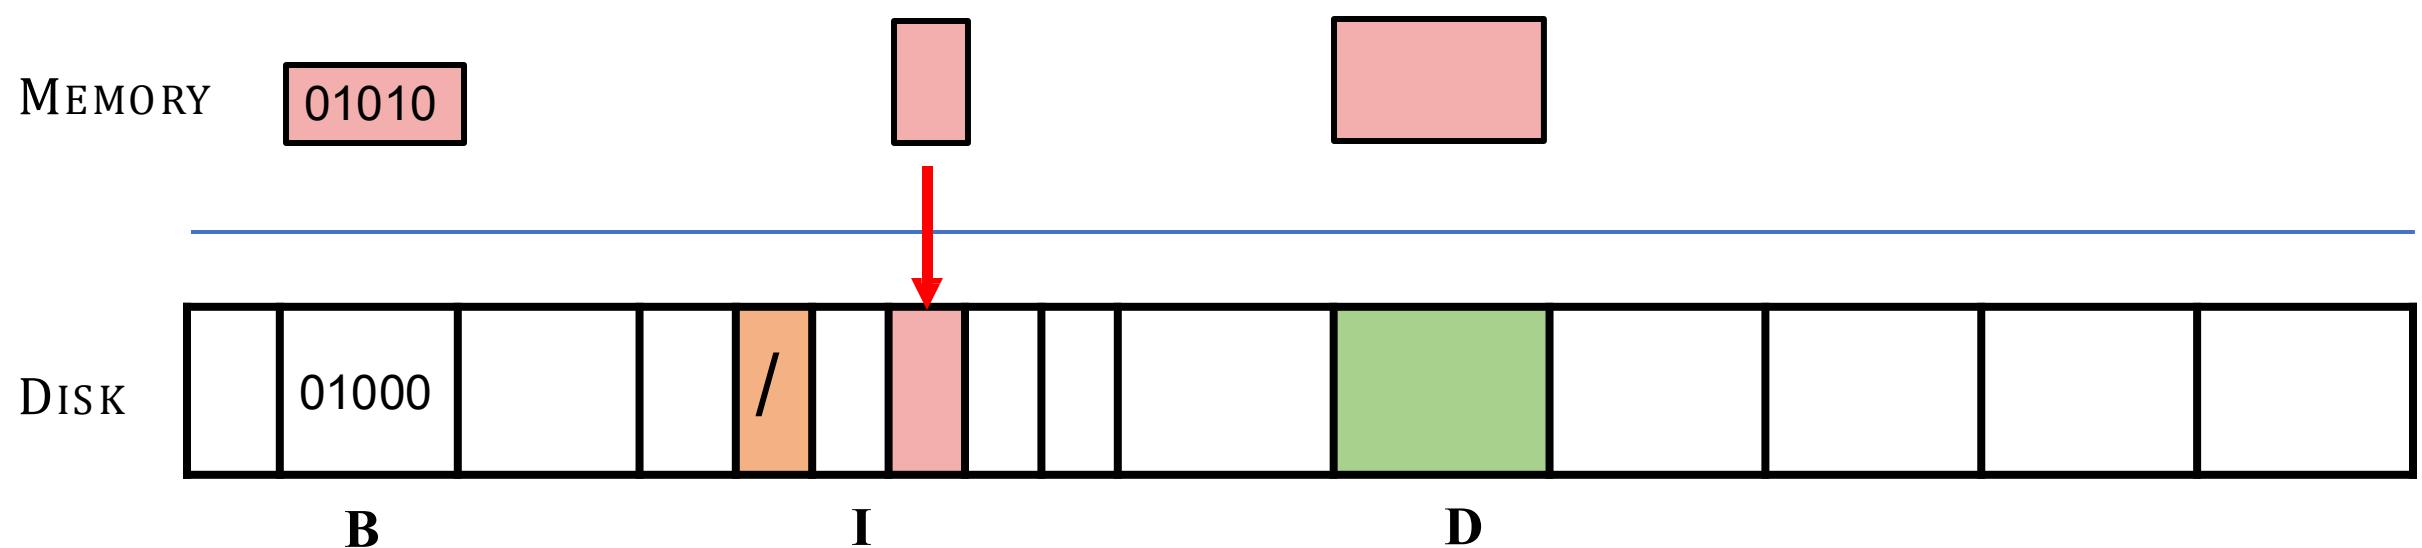


Example: Inode First

Write Ordering: Inode (I), Bitmap (B), Data (D)

- But CRASH after I has reached disk, before B or D

Result?

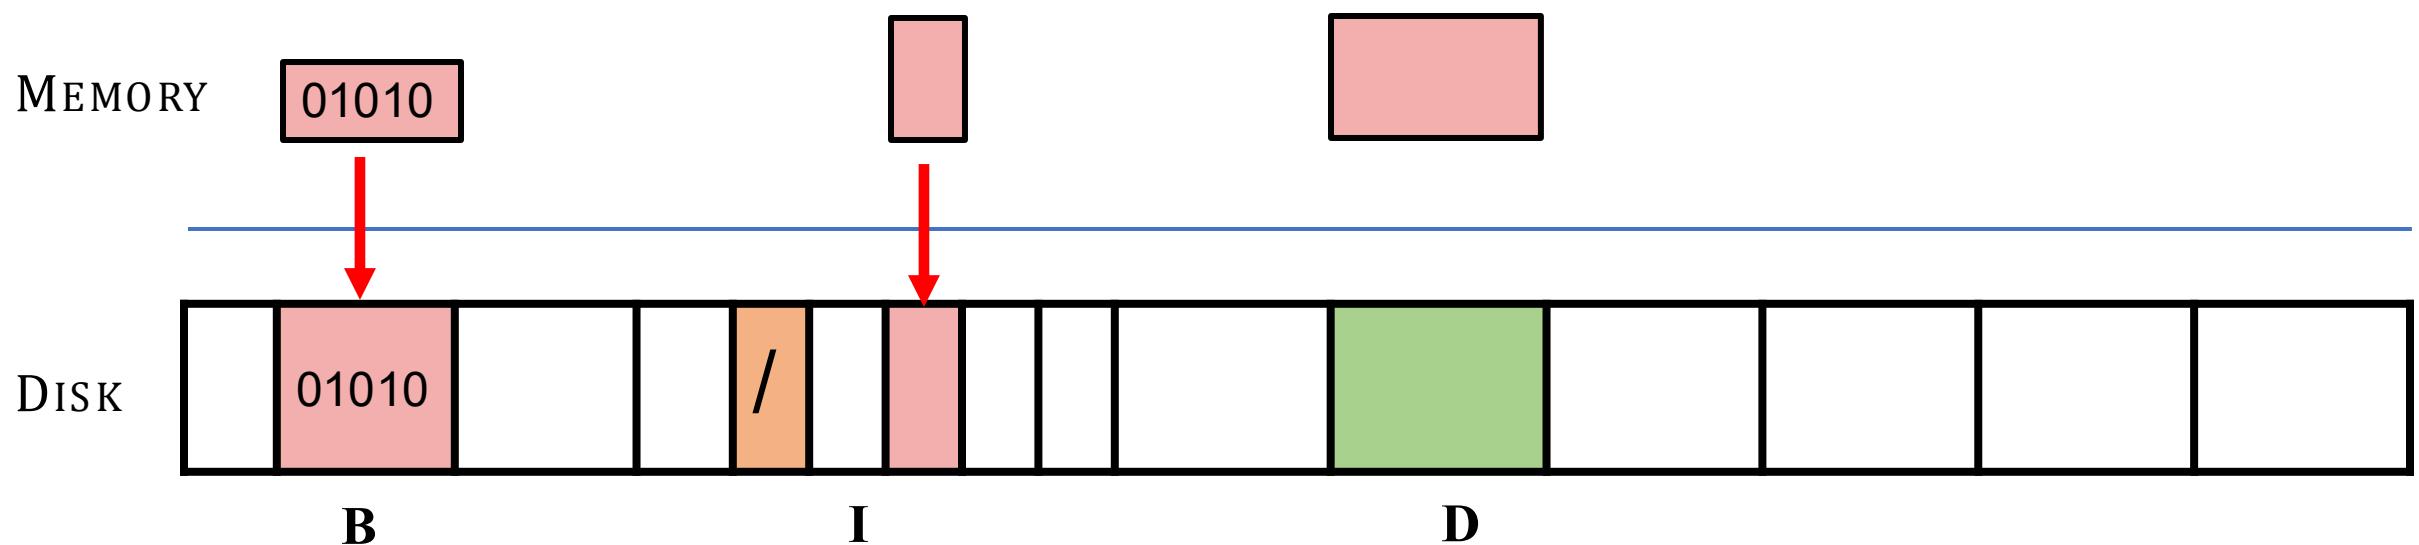


Example: Inode First

Write Ordering: Inode (I), Bitmap (B), Data (D)

- But CRASH after I AND B have reached disk, before D

Result?



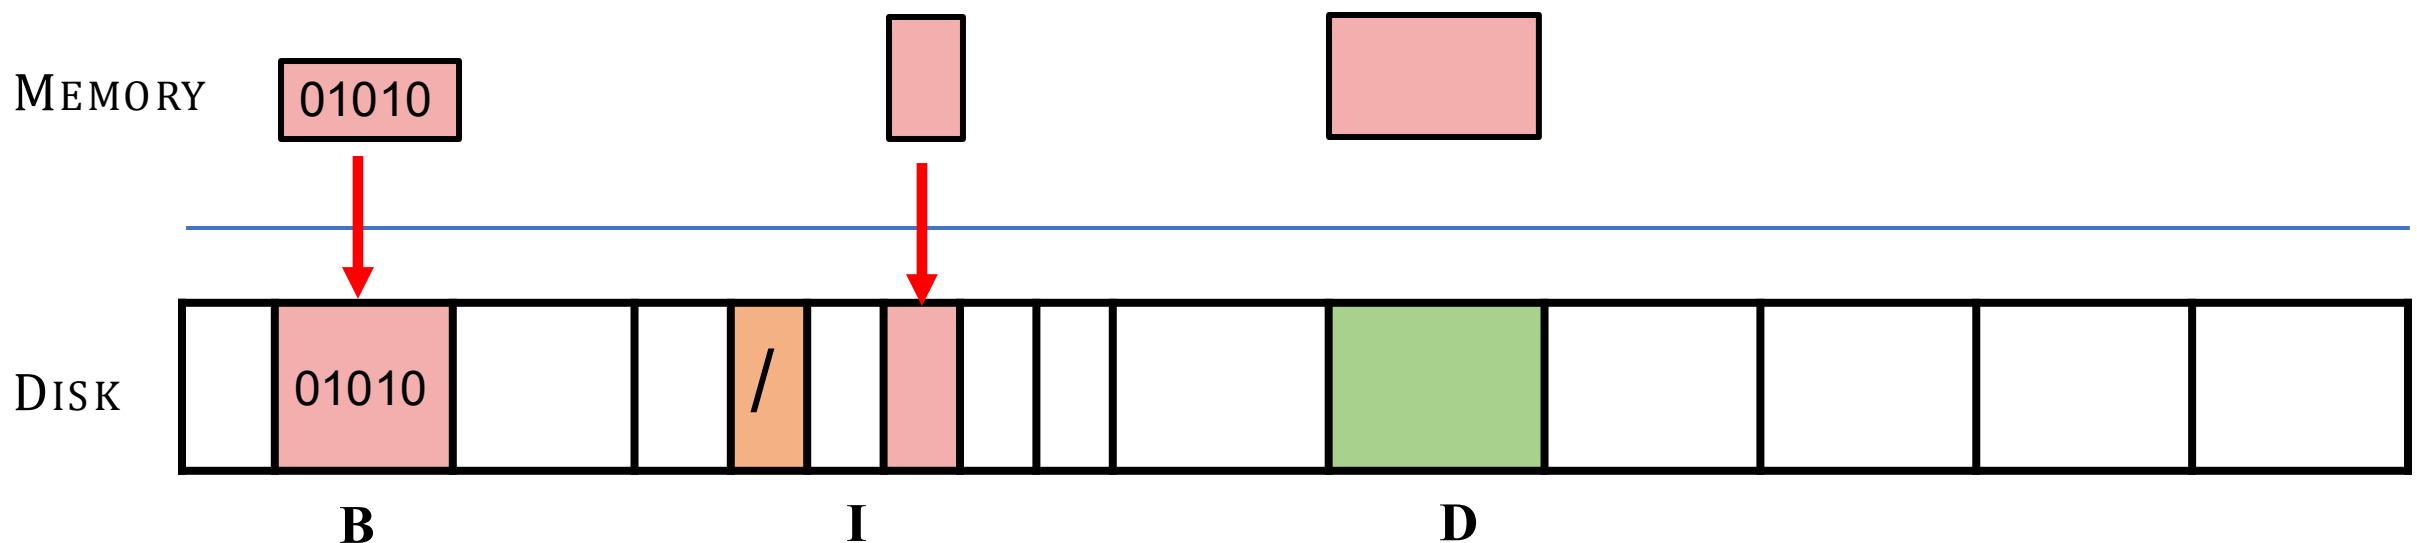
Example: Inode First

Write Ordering: Inode (I), Bitmap (B), Data (D)

- But CRASH after I AND B have reached disk, before D

Result?

- What if data block is a new block for the new file (i.e., create file with data)

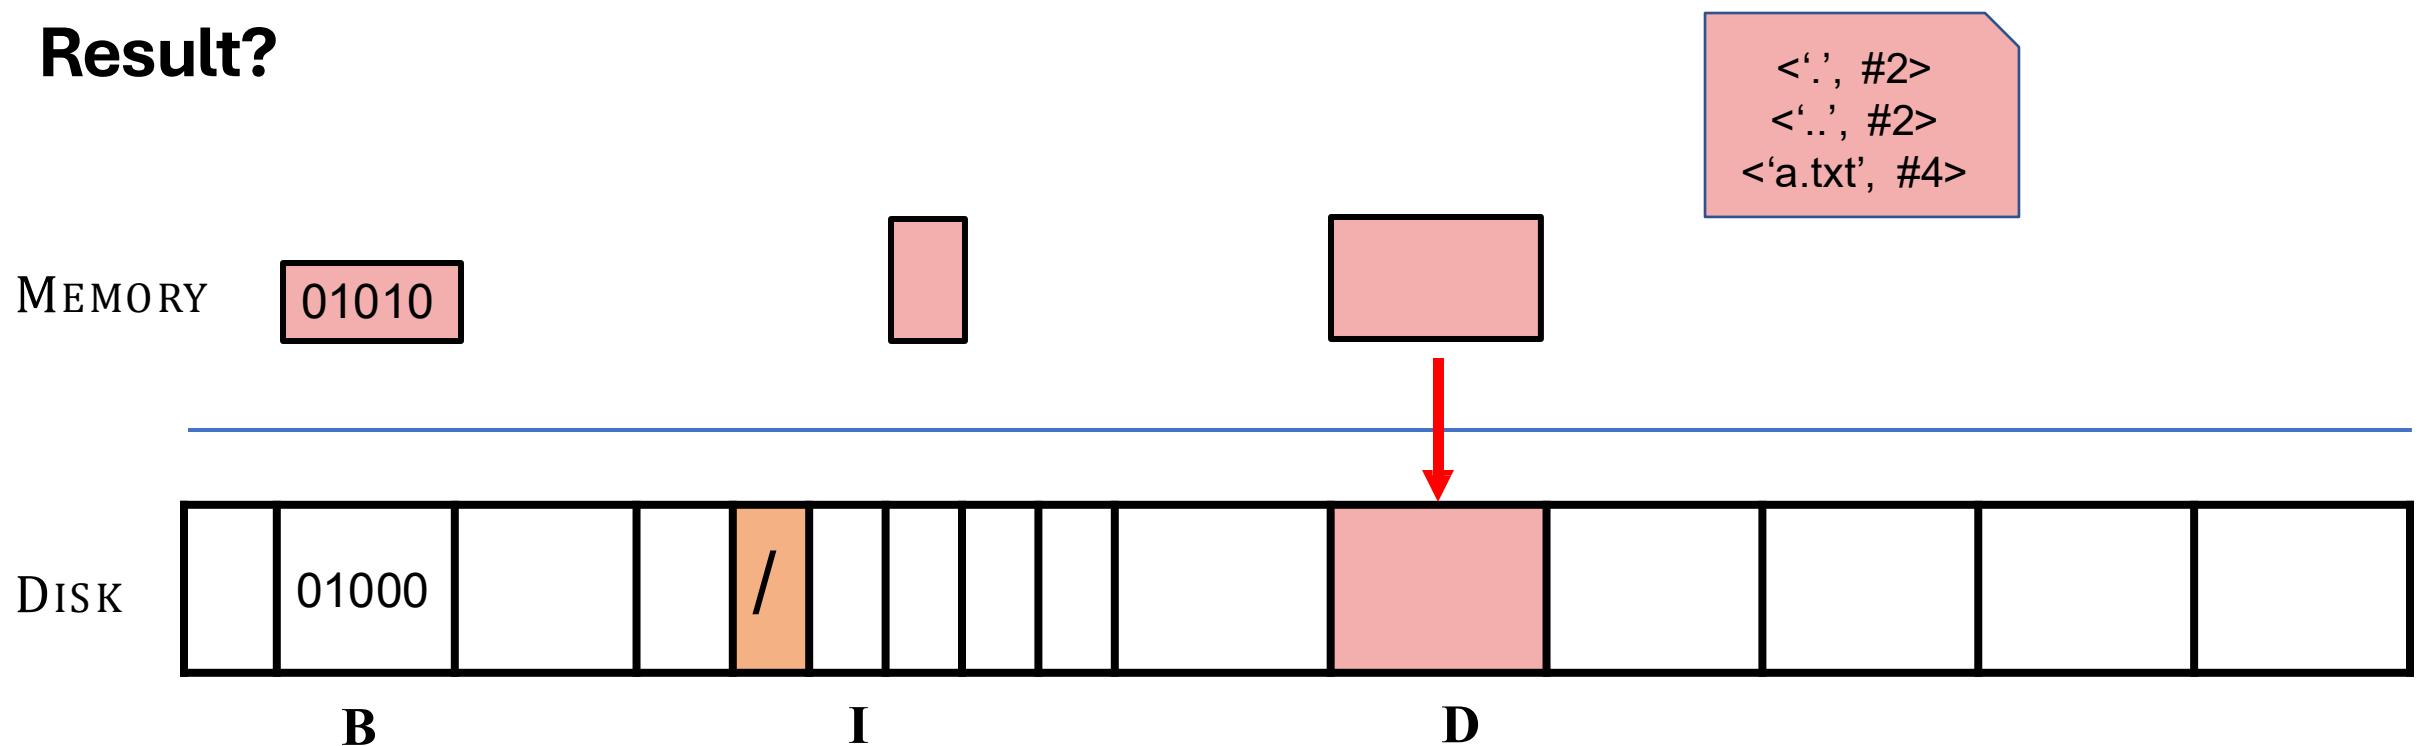


Example: Data First

Write Ordering: Data (D) , Bitmap (B), Inode (I)

- CRASH after D has reached disk, before I or B

Result?



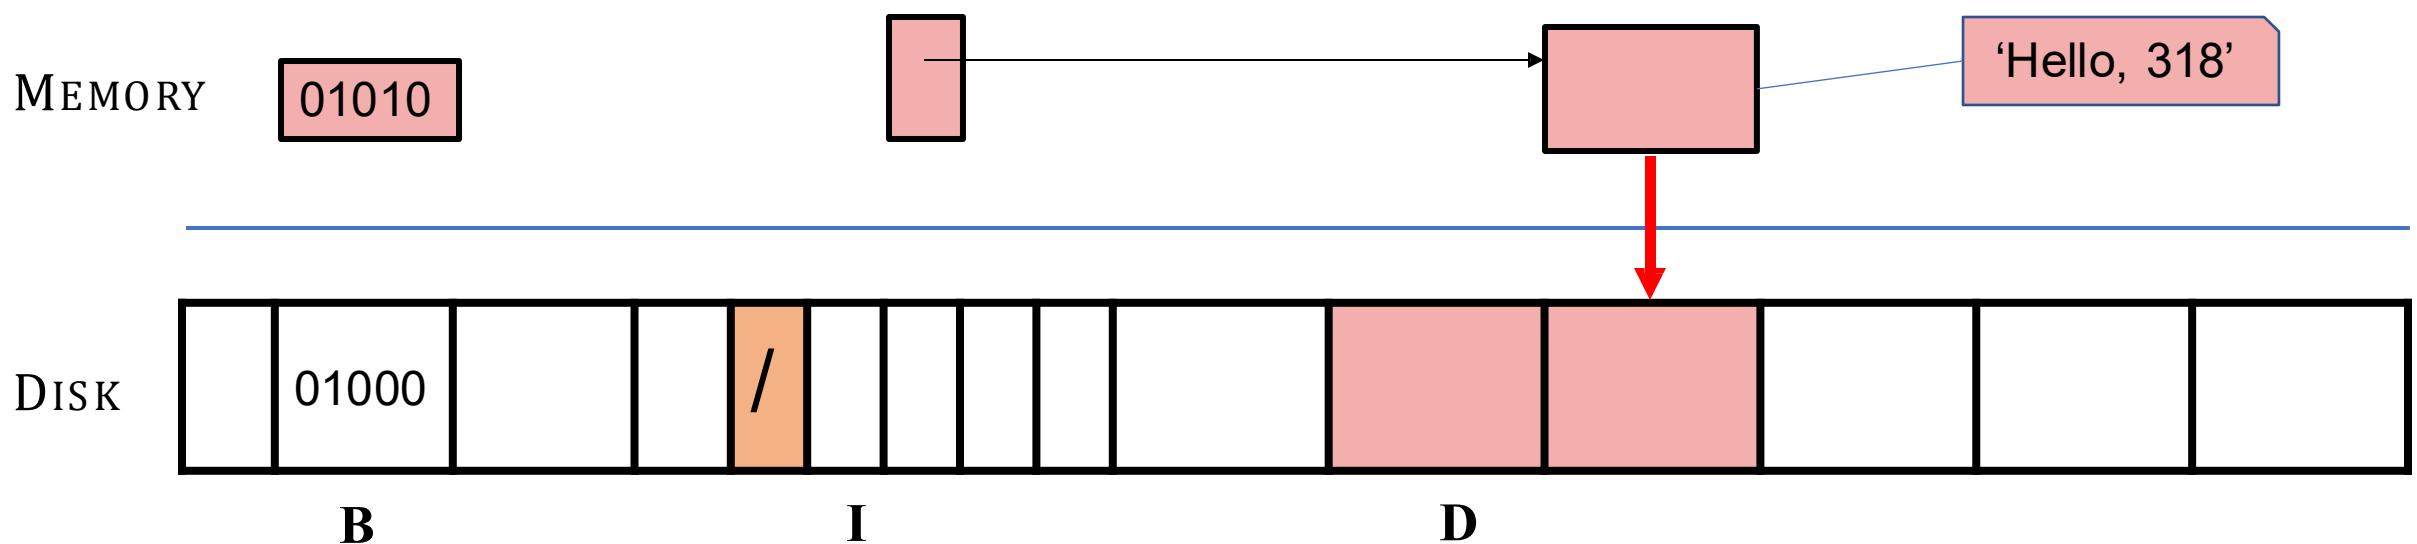
Example: Data First

Write Ordering: Data (D) , Bitmap (B), Inode (I)

- CRASH after D has reached disk, before I or B

Result?

- What if data block is a new block for the new file (i.e., create file with data)



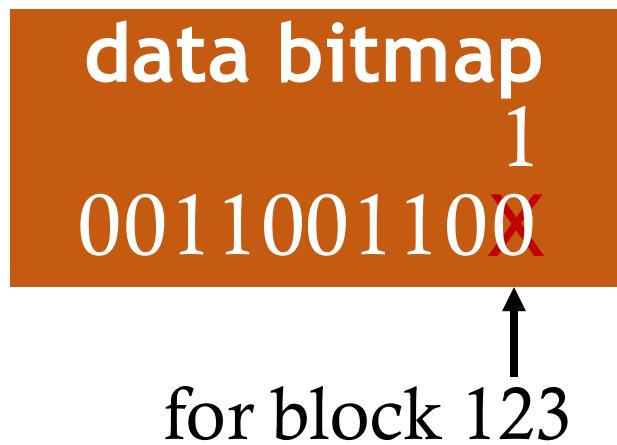
Traditional Solution: fsck

FSCK: “file system checker”

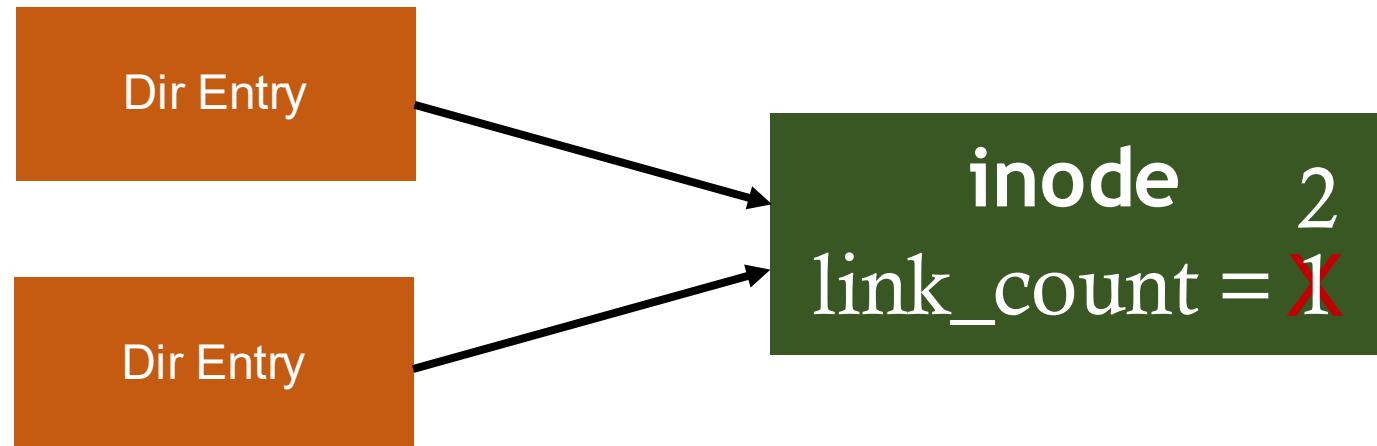
When system boots:

- Make multiple passes over file system, looking for inconsistencies
 - e.g., inode pointers and bitmaps, directory entries and inode reference counts
- Try to fix automatically

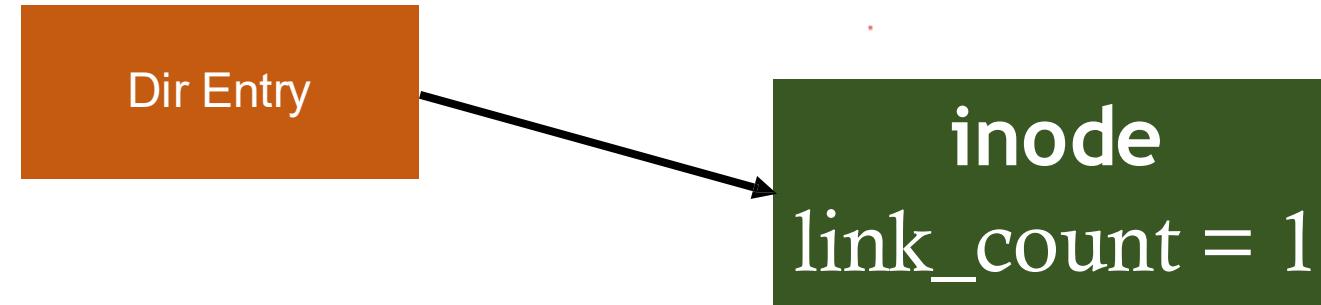
FSCK Example 1



FSCK Example 2

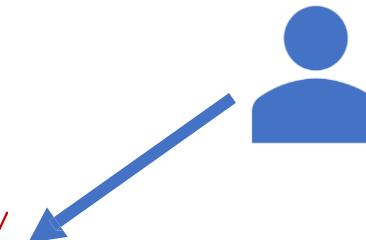


FSCK Example 3

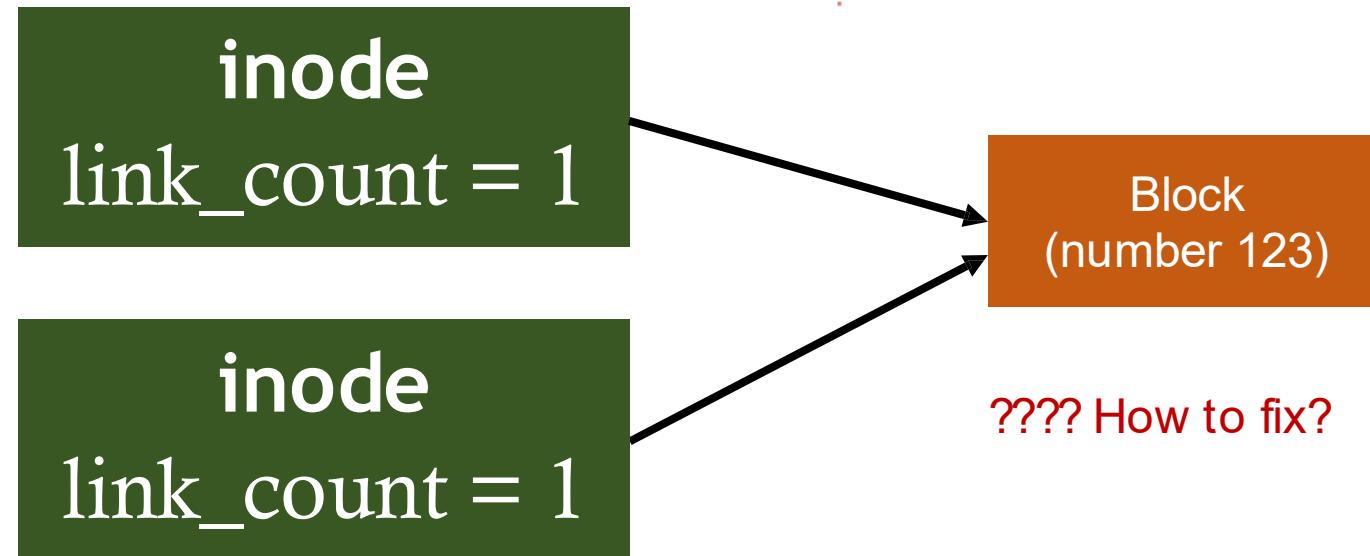


???? How to fix?

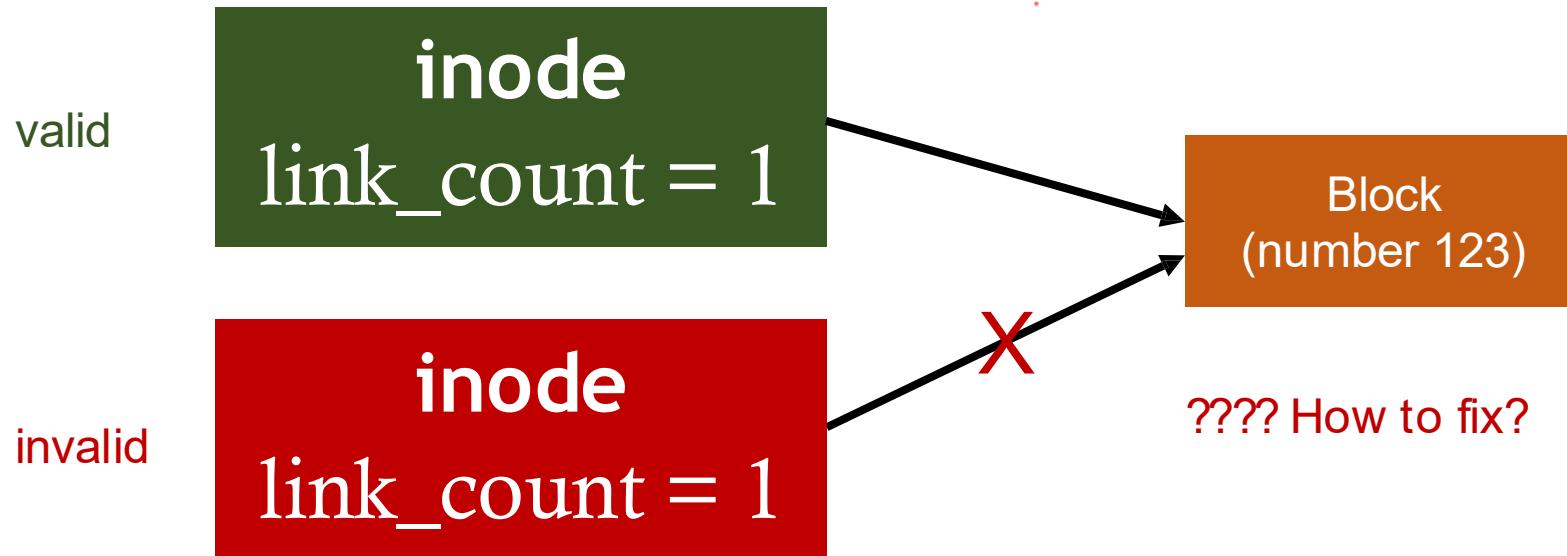
```
ls -l /
total 150
drwxr-xr-x  401 18432 Dec 31 1969 afs/
drwxr-xr-x.  2 4096 Nov  3 09:42 bin/
drwxr-xr-x.  5 4096 Aug  1 14:21 boot/
dr-xr-xr-x.  13 4096 Nov  3 09:41 lib/
dr-xr-xr-x.  10 12288 Nov  3 09:41 lib64/
drwx-----.  2 16384 Aug  1 10:57 lost+found/
...
.
```



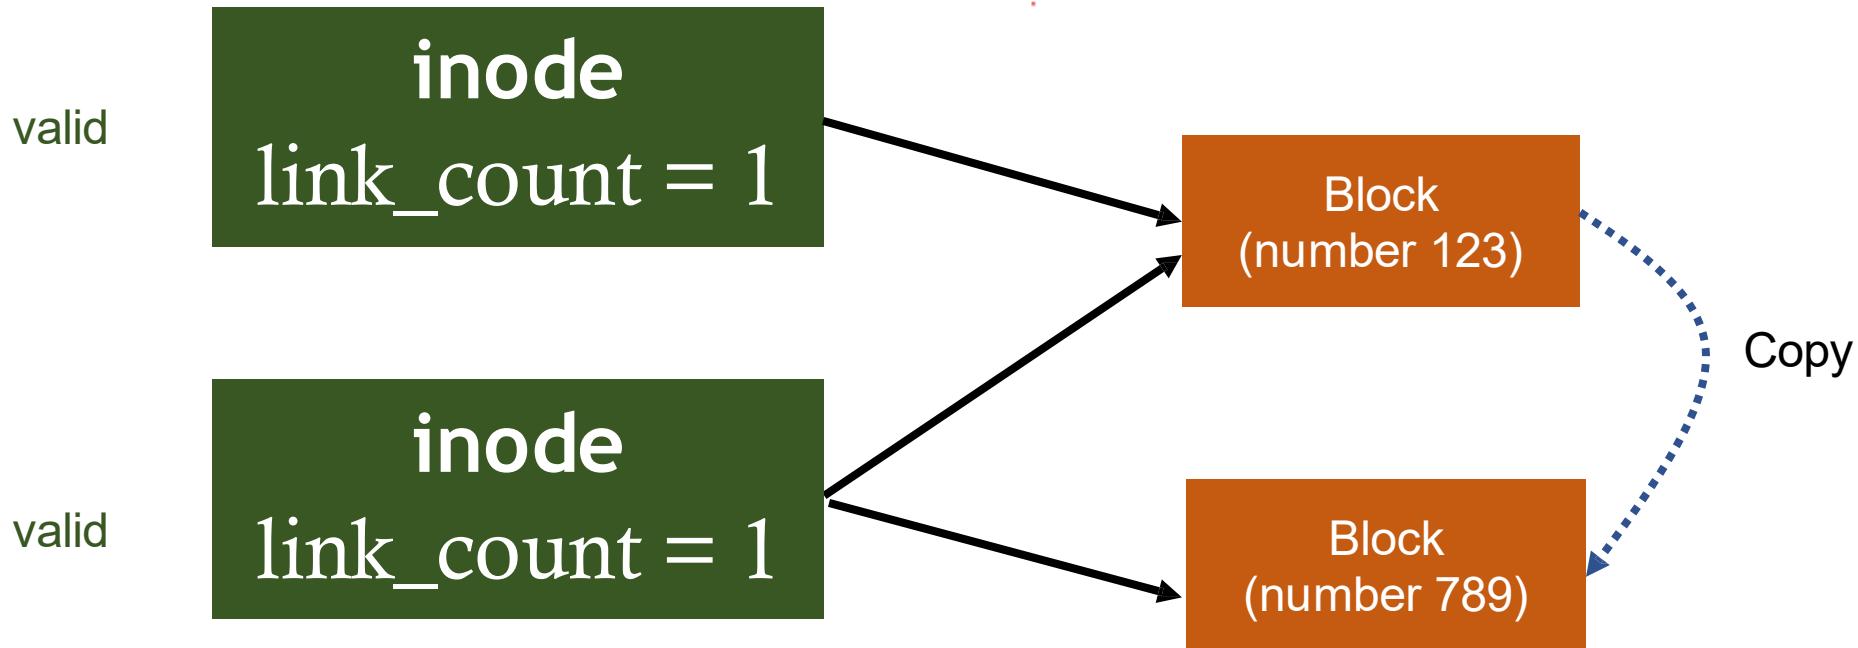
FSCK Example 4



FSCK Example 4.a



FSCK Example 4.b



Traditional Solution: FSCK

FSCK: “file system checker”

When system boots:

- Make multiple passes over file system, looking for inconsistencies
- Try to fix automatically
 - Example: **B' I D, B I' D**
- Or punt to admin
- Check lost+found, manually put the missing-link files to the correct place

Traditional Solution: FSCK

Problem:

- Cannot fix all crash scenarios
 - Can B' | D' be fixed?
- Performance
 - Sometimes takes hours to run
 - Checking a 600GB disk takes ~70 minutes
 - Does fsck have to run upon every reboot?
- Not well-defined consistency

Another Solution: Journaling

Idea: Write “intent” down to disk before updating file system

- Called the “[Write Ahead Logging](#)” or “journal”
- Originated from database community

When crash occurs, look through log to see what was going on

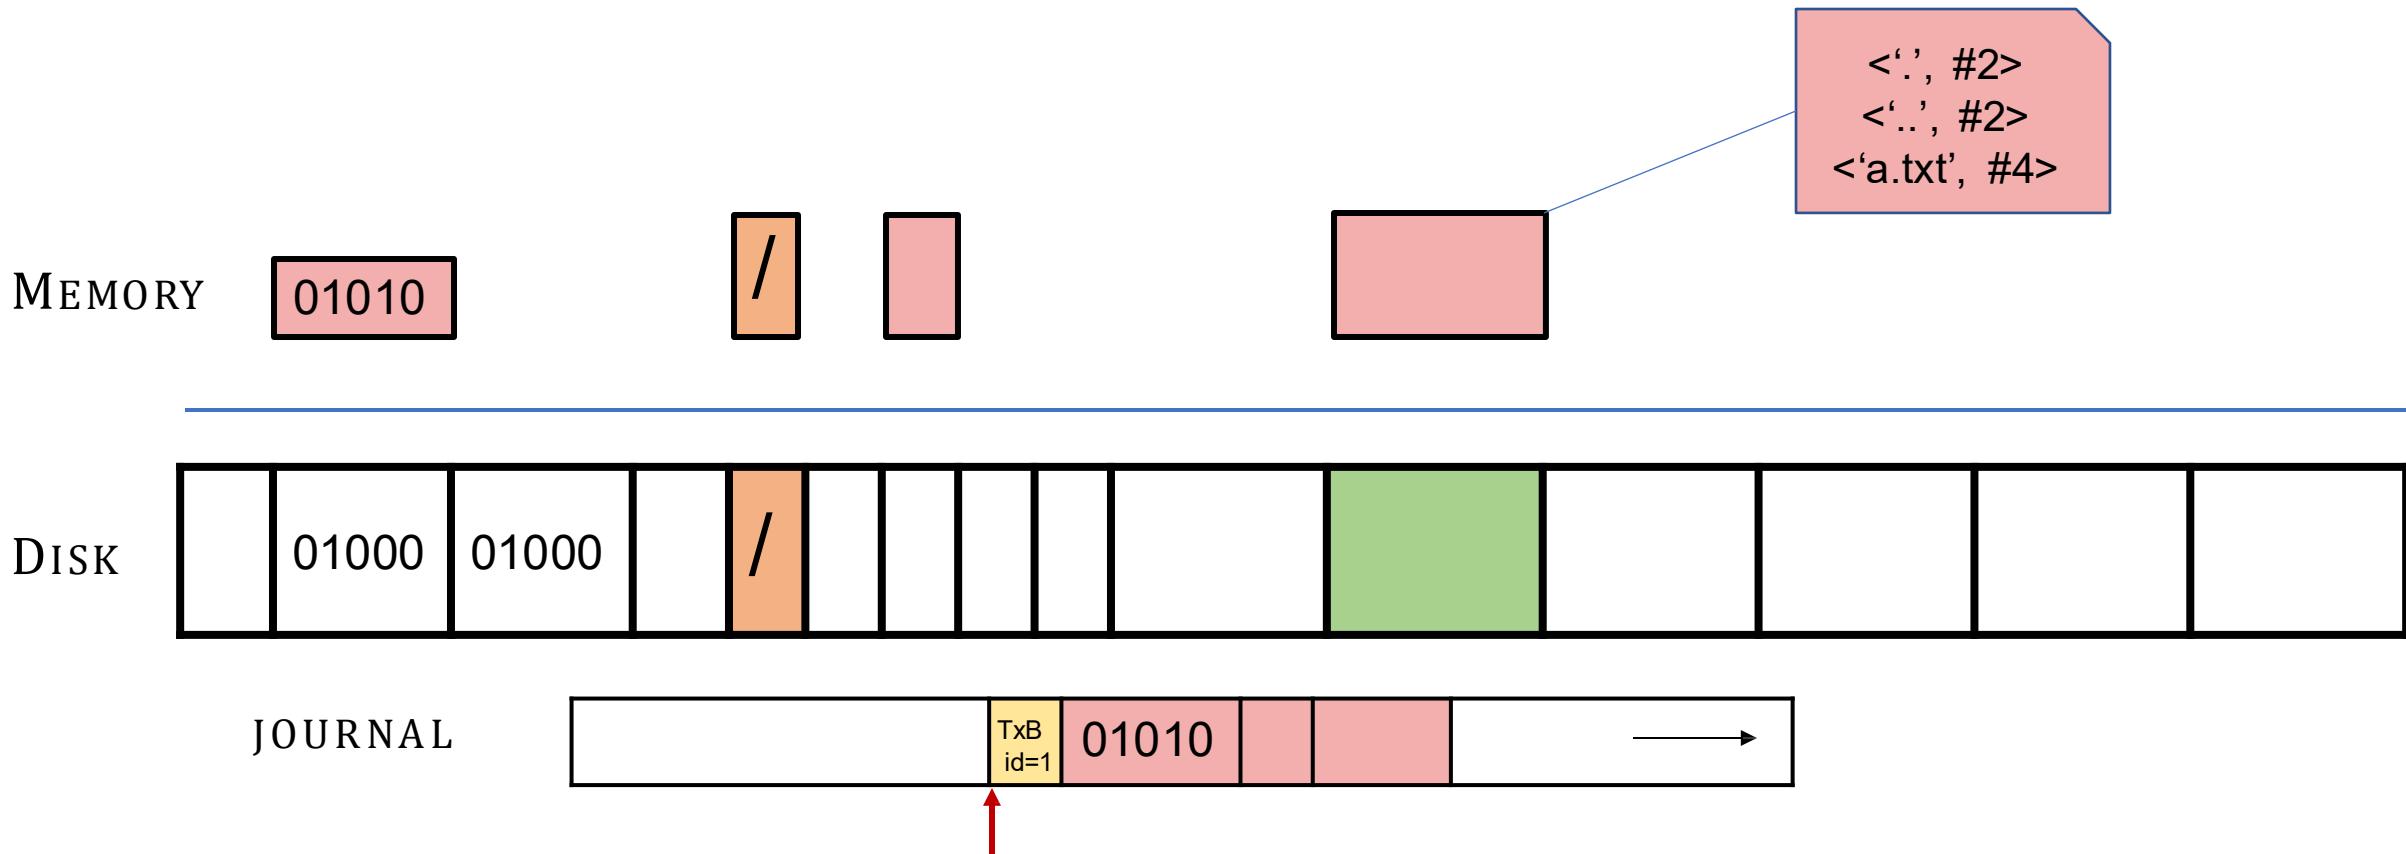
- Use contents of log to fix file system structures
 - Crash before “intent” is written → no-op
 - Crash after “intent” is written → redo op
- The process is called “recovery”

Case Study: Linux Ext3

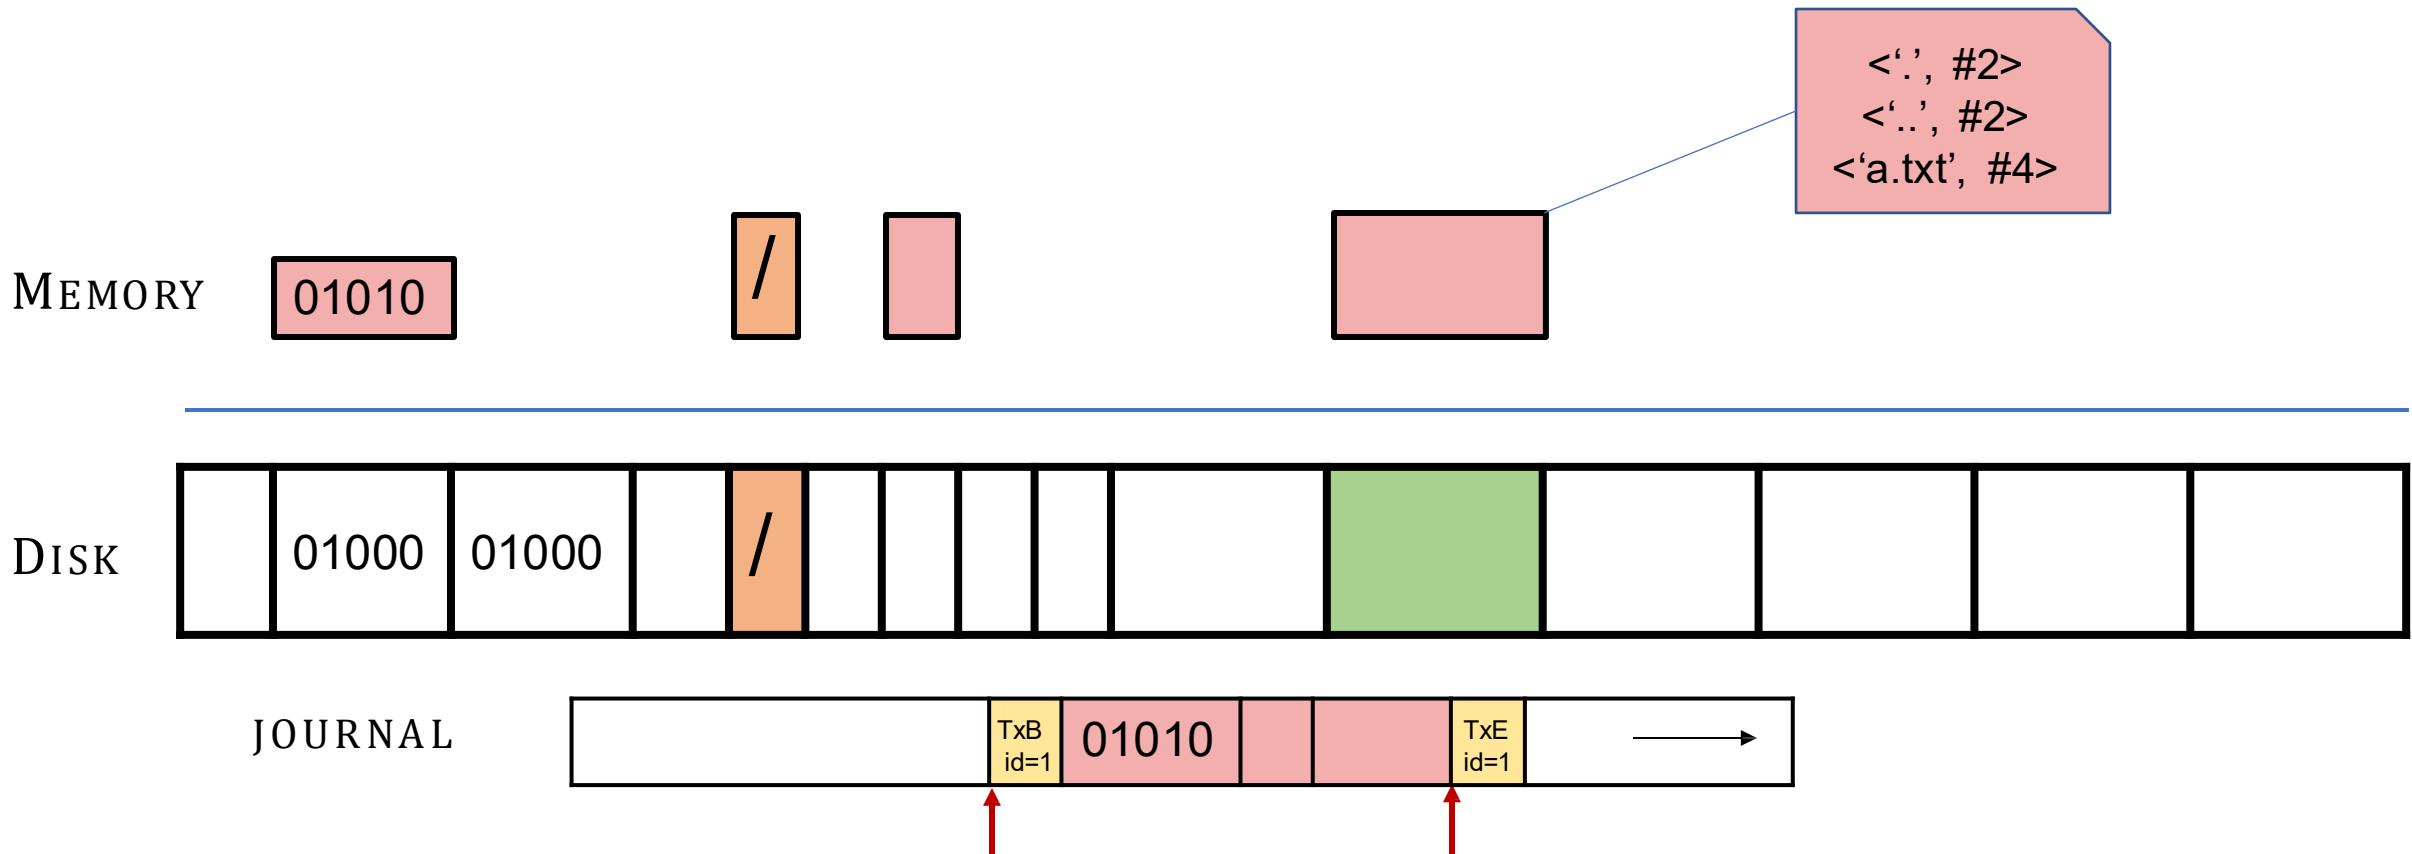
Write real block contents of the update to log

- Four totally ordered steps:
 1. Commit dirty blocks to journal as one transaction (TxBegin, I, B, D blocks)
 2. Write commit record (TxEnd)
 3. Copy dirty blocks to real file system (**checkpointing**)
 4. Reclaim the journal space for the transaction

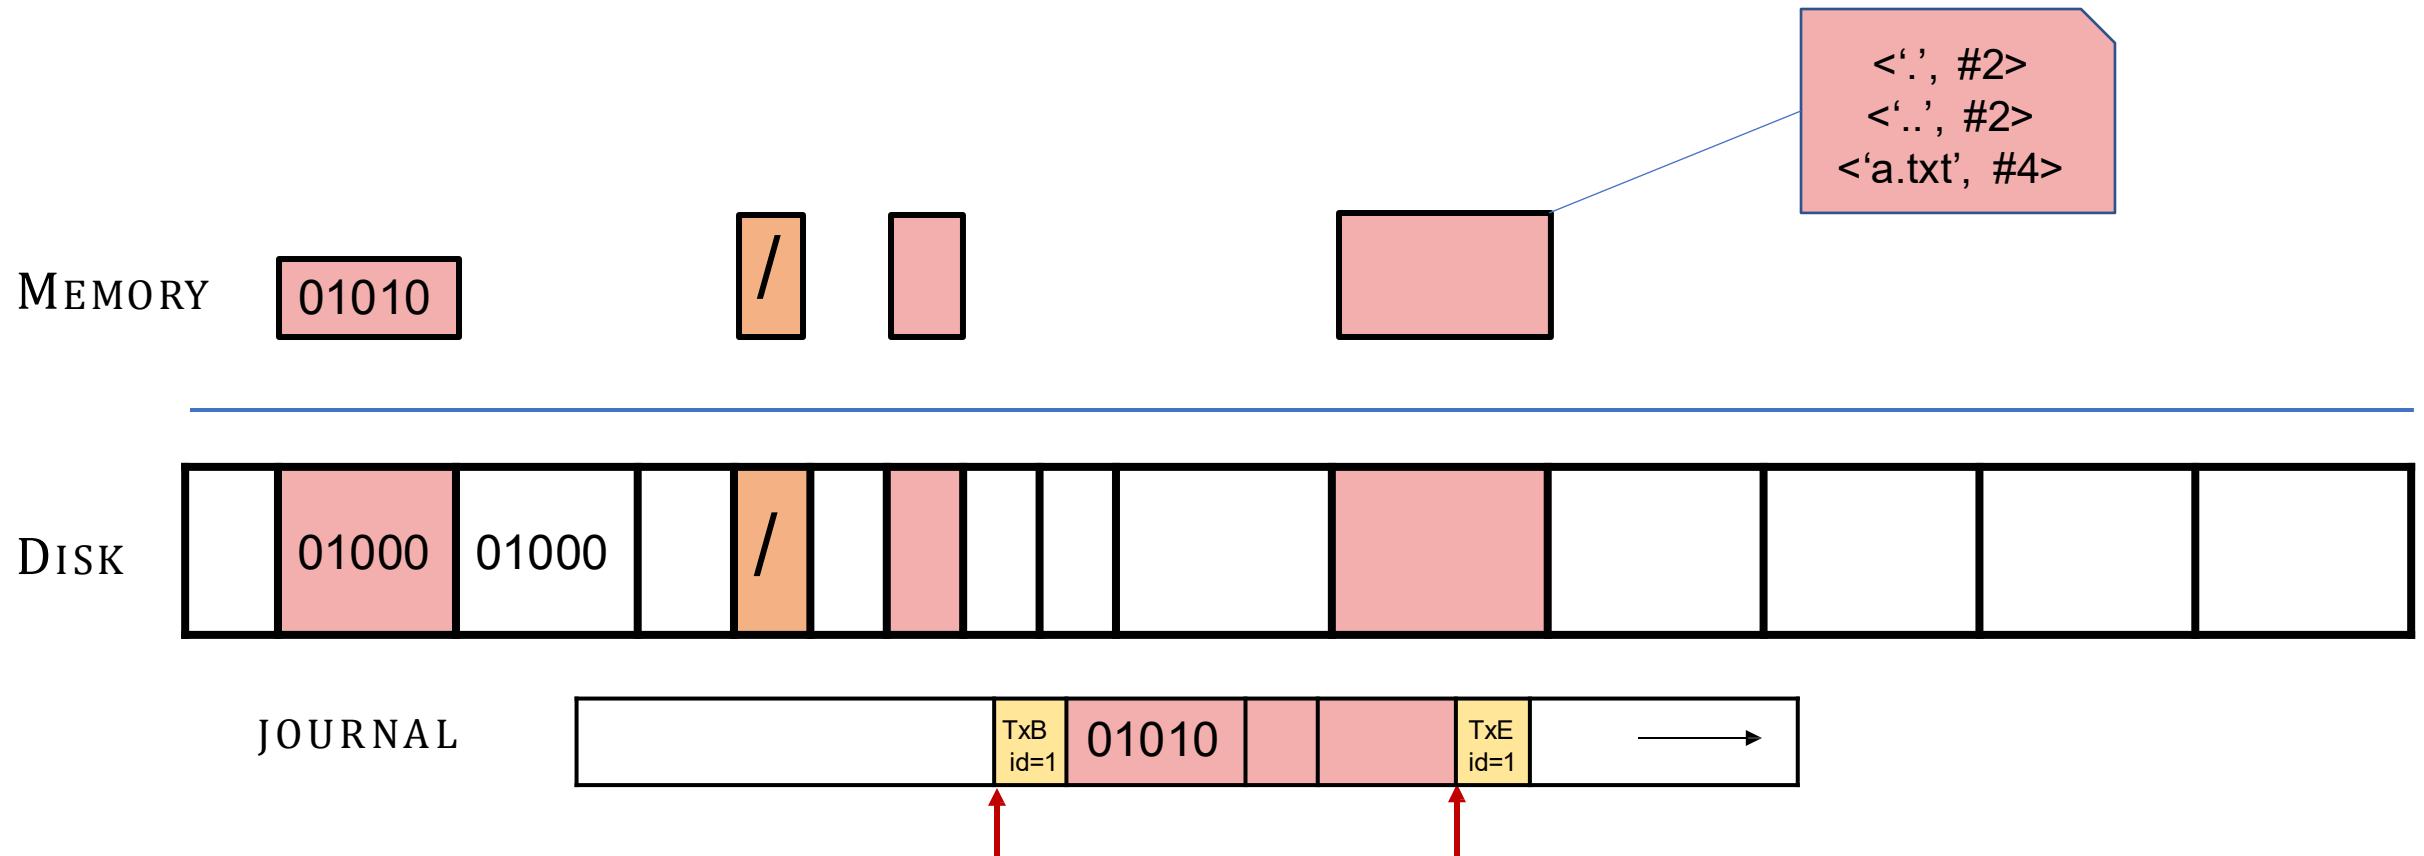
Step 1. Write Blocks to Journal



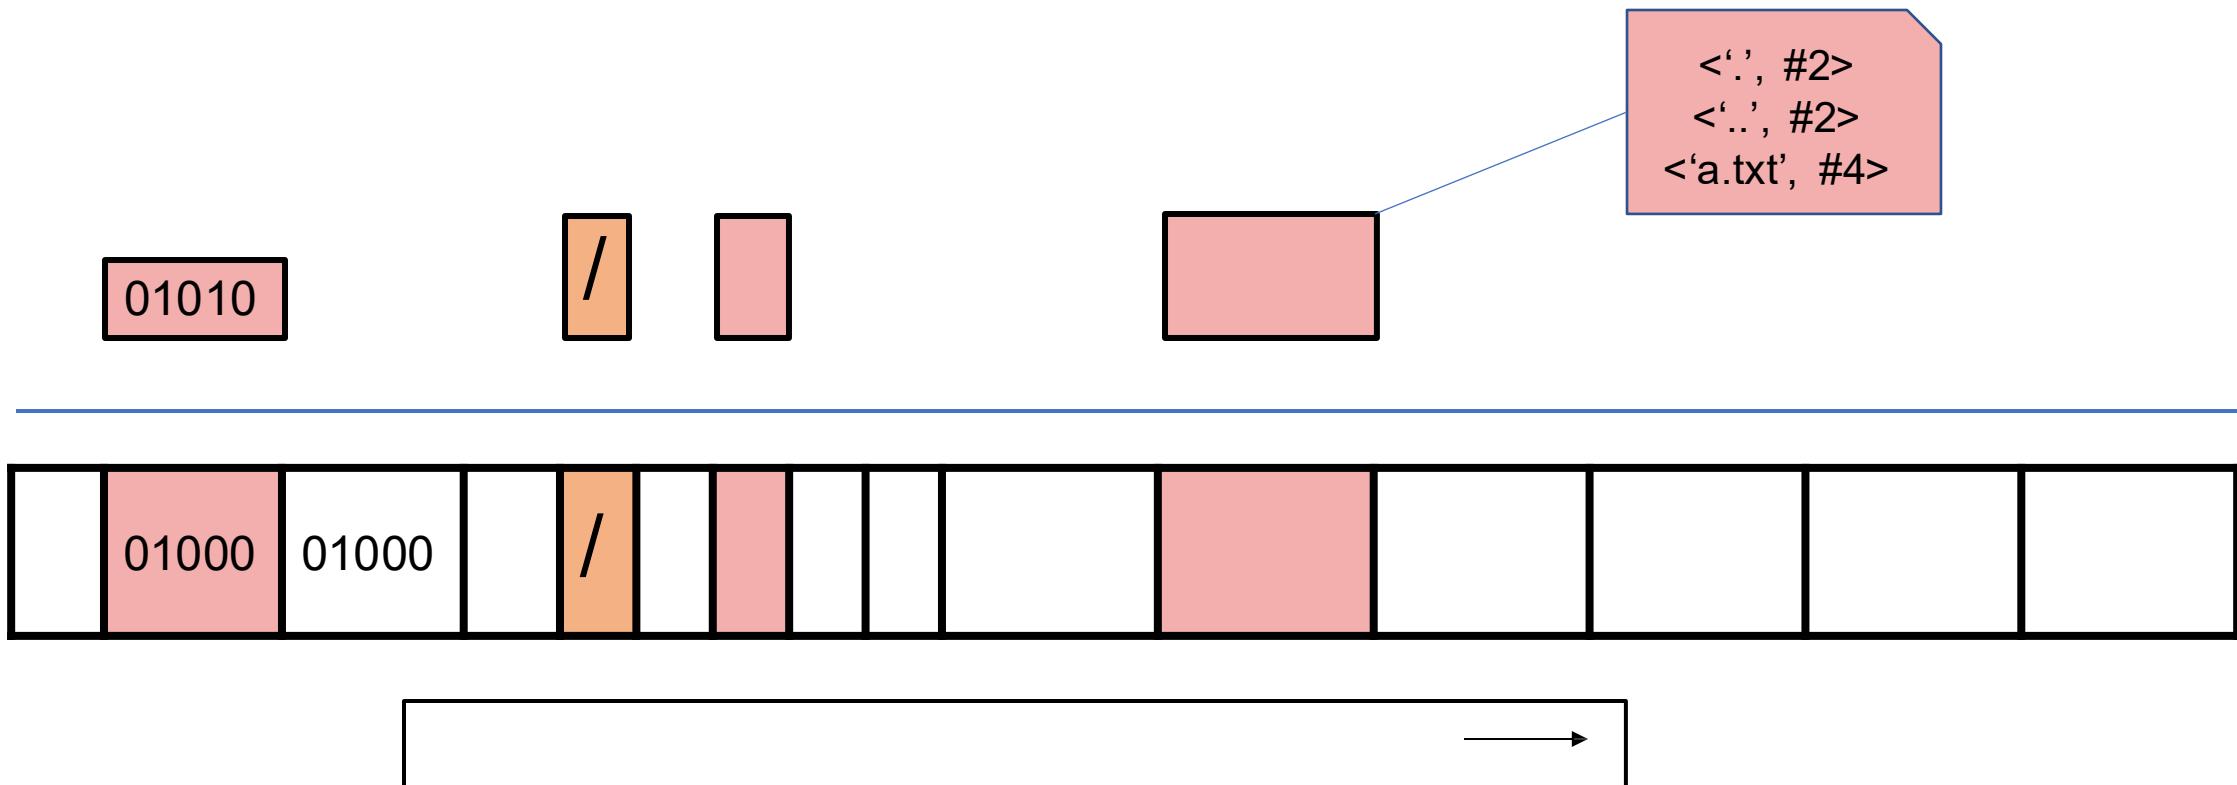
Step 2. Write Commit Record



Step 3: Copy Dirty Blocks to Real FS



Step 4: Reclaim Journal Space

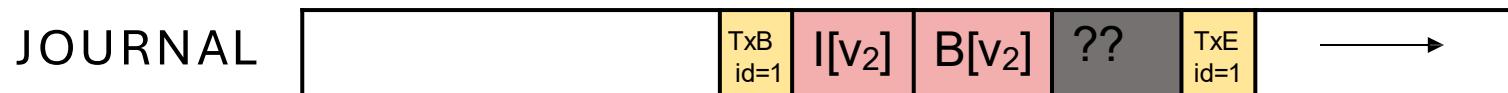


What If There Is A Crash?

Recovery: Go through log and “redo” operations that have been successfully committed to log

What if ...

- TxBegin but not TxEnd in log?
- TxBegin through TxEnd are in log, but D has not reached the journal?



- How could this happen?
- Why don't we merge step 2 and step 1?
- Tx in log, I, B, D have been checkpointed, but Tx is not freed from log?

Summary of Journaling Write Orders

Journal writes < FS writes

- Otherwise, crash → FS broken, but no record in journal to patch it up

FS writes < Journal clear

- Otherwise, crash → FS broken, but record in journal is already cleared

Journal writes < commit record write < FS writes

- Otherwise, crash → record appears committed, but contains garbage

Ext3 Journaling Modes

Journaling has cost

- one write = two disk writes, two seeks

Several journaling modes balance consistency and performance

Data journaling: journal all writes, including file data

- Problem: expensive to journal data

Metadata journaling: journal only metadata

- Used by most FS (IBM JFS, SGI XFS, NTFS)
- Problem: file may contain garbage data

Ordered mode: write file data to real FS first, then journal metadata

- Default mode for ext3
- Problem: old file may contain new data

Summary

The consistent update problem

- Example of file creation and different crash scenarios

Two approaches to crash consistency

- FSCK: slow, not well-defined consistency
- Journaling: well-defined consistency, different modes

Other approach

- Soft updates (advanced OS topics)